

# **Convolutional Neural Networks**

**Shusen Wang**

# Inner Products

内积

# Vector Inner Product

- $\mathbf{a} = \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix}$  and  $\mathbf{b} = \begin{bmatrix} 5 \\ 6 \\ 7 \\ 8 \end{bmatrix}$ .

- Inner product:

$$\langle \mathbf{a}, \mathbf{b} \rangle = \mathbf{a}^T \mathbf{b} = \sum_i a_i b_i = 70.$$

# Matrix Inner Product

- $\mathbf{A} = \begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$  and  $\mathbf{B} = \begin{bmatrix} 5 & 7 \\ 6 & 8 \end{bmatrix}$ .
- Inner product:  
$$\langle \mathbf{A}, \mathbf{B} \rangle = \sum_i \sum_j a_{ij} b_{ij} = 70.$$
- Property:  $\langle \mathbf{A}, \mathbf{B} \rangle = \langle \text{vec}(\mathbf{A}), \text{vec}(\mathbf{B}) \rangle$ .

# Tensor Inner Product

- $\mathbf{A} = \begin{matrix} & \\ & \text{[3x3]} \\ \text{[3x3]} & \end{matrix}$  and  $\mathbf{B} = \begin{matrix} & \\ & \text{[3x3]} \\ \text{[3x3]} & \end{matrix}.$

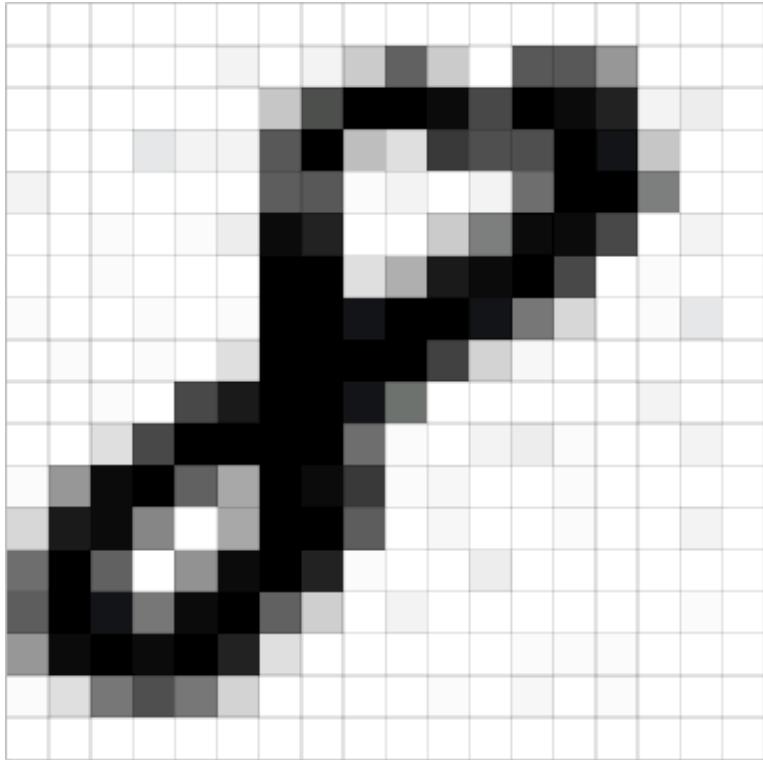
- Inner product:

$$\langle \mathbf{A}, \mathbf{B} \rangle = \sum_i \sum_j \sum_k a_{ijk} b_{ijk}.$$

- Property:  $\langle \mathbf{A}, \mathbf{B} \rangle = \langle \text{vec}(\mathbf{A}), \text{vec}(\mathbf{B}) \rangle.$

# Convolution for Images

# An Image is A Matrix/Tensor of Pixel Values



## The MNIST Dataset

- Each grayscale image,  $\mathbf{x}_j$ , is a  $28 \times 28$  matrix.
- (Previously, we reshape it to a 784-dim vector.)

# Convolution

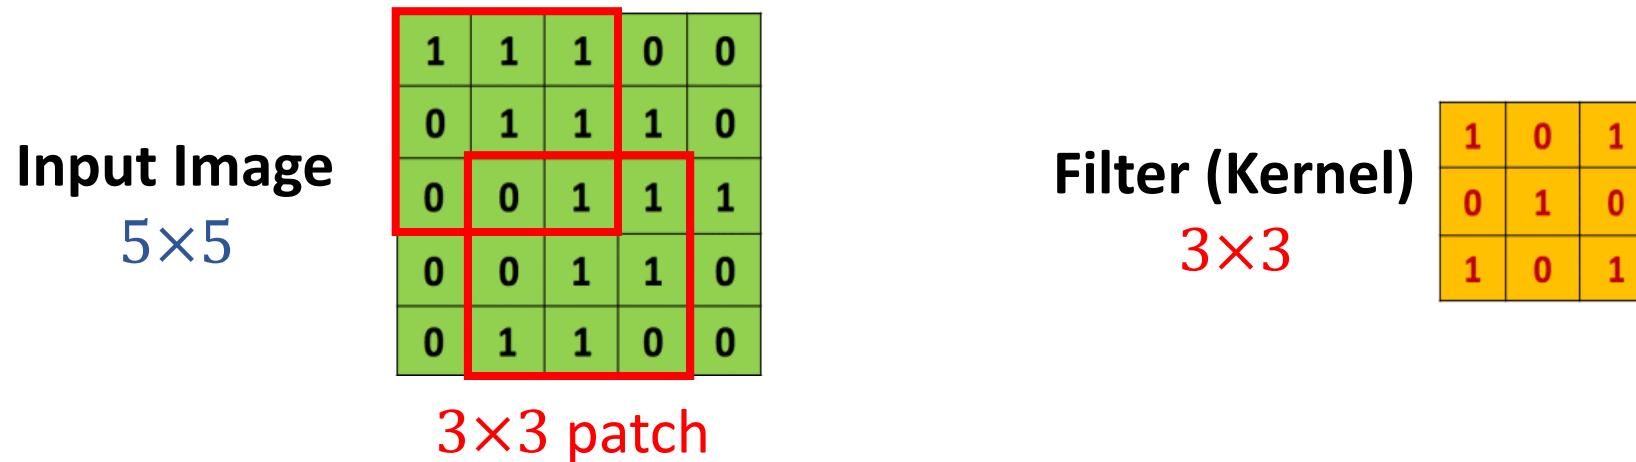
**Input Image**  
 $5 \times 5$

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

**Filter (Kernel)**  
 $3 \times 3$

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

# Convolution



**Question:** how many  $3 \times 3$  patches does a  $5 \times 5$  image have?

**Answer:**  $(5 - 3 + 1) \times (5 - 3 + 1) = 9.$

# Convolution

Input Image  
 $5 \times 5$

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Filter (Kernel)  
 $3 \times 3$

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |



Convolution:

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

Result  
 $3 \times 3$

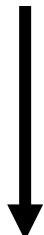
# Convolution

Input Image  
 $5 \times 5$

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Filter (Kernel)  
 $3 \times 3$

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |



Convolution:

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

Result  
 $3 \times 3$

The value **4** is the inner product of the patch  
内积

|   |   |   |
|---|---|---|
| 1 | 1 | 1 |
| 0 | 1 | 1 |
| 0 | 0 | 1 |

and the filter

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

# Convolution

Input Image  
 $5 \times 5$

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Filter (Kernel)  
 $3 \times 3$

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

Convolution:

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

Result  
 $3 \times 3$

The value **3** is the inner product of the patch  
内积

|   |   |   |
|---|---|---|
| 1 | 1 | 0 |
| 1 | 1 | 1 |
| 0 | 1 | 1 |

and the filter

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

# Convolution

Input Image  
 $5 \times 5$

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Filter (Kernel)  
 $3 \times 3$

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |



Convolution:

|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

Result  
 $3 \times 3$

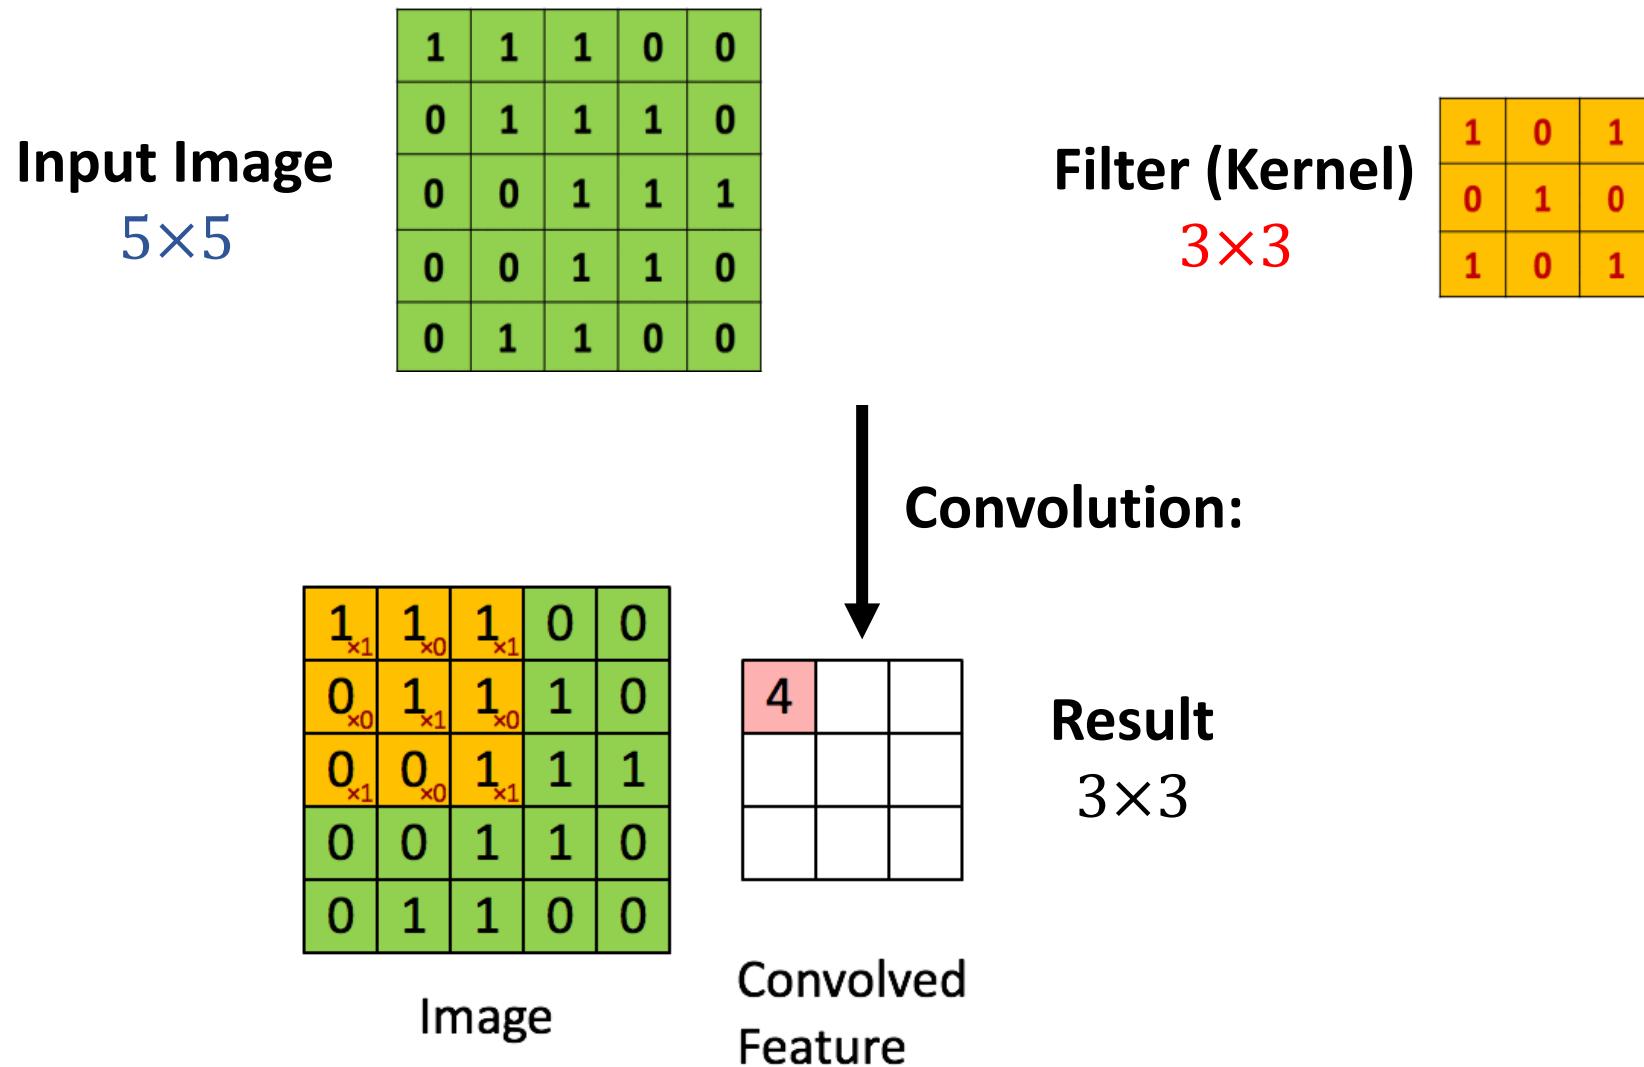
The value **4** is the inner product of the patch

|   |   |   |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 0 |
| 1 | 0 | 0 |

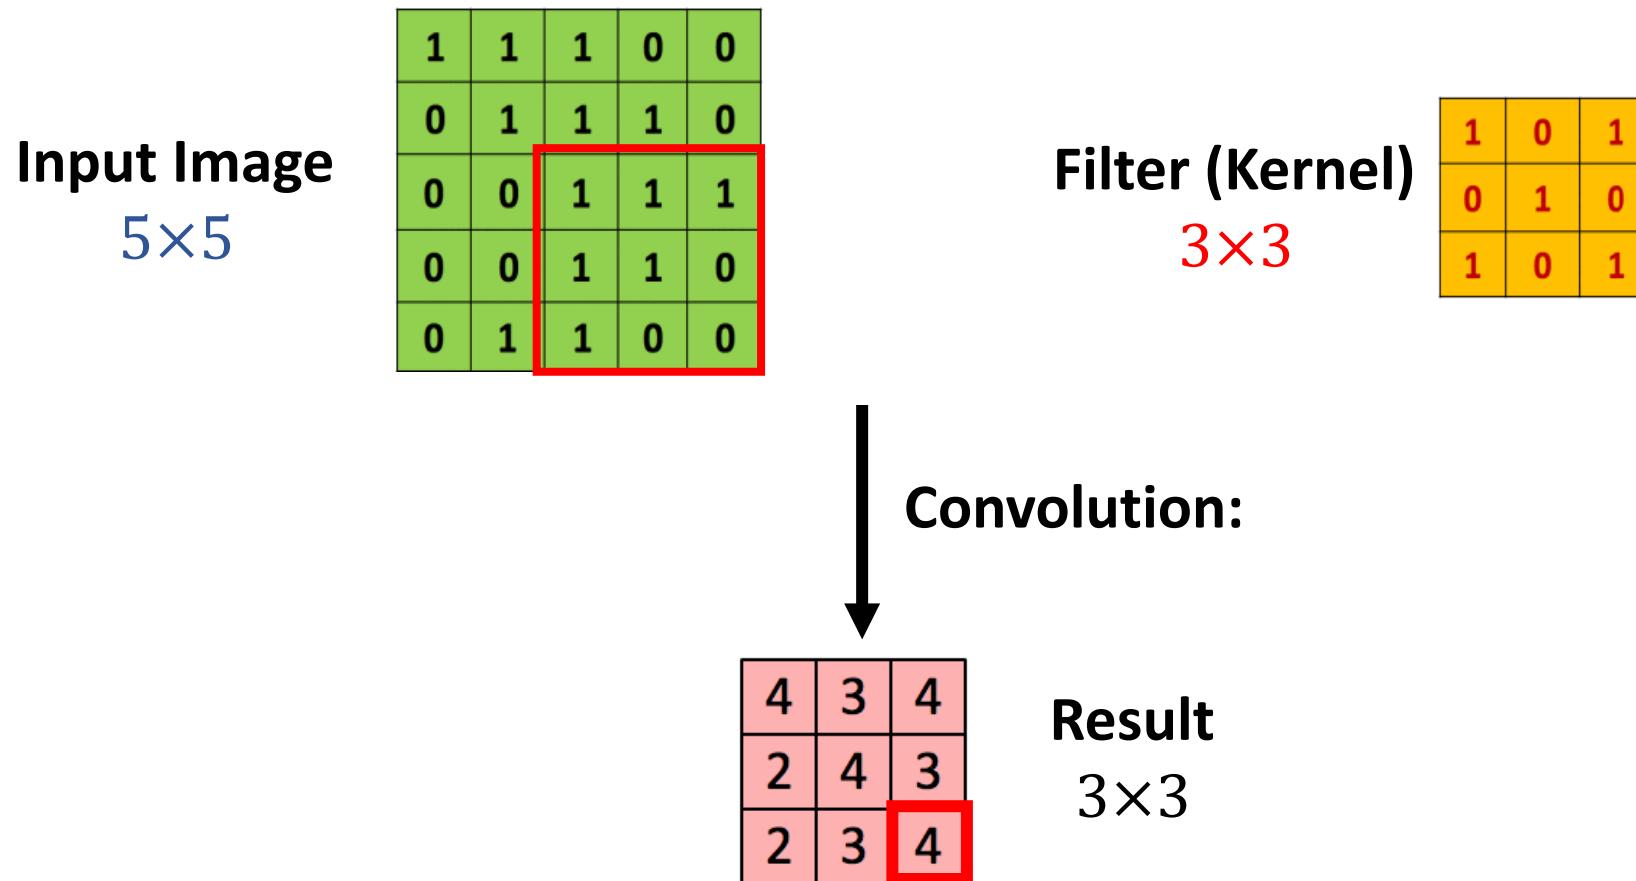
and the filter

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

# Convolution



# Convolution



**Question:** Why is the result a  $3 \times 3$  matrix?

# Convolution

**Input Image**  
 $4 \times 4$

|   |   |   |   |
|---|---|---|---|
| 1 | 1 | 1 | 0 |
| 0 | 1 | 1 | 1 |
| 0 | 1 | 1 | 0 |
| 1 | 1 | 0 | 0 |

**Filter (Kernel)**  
 $3 \times 3$

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

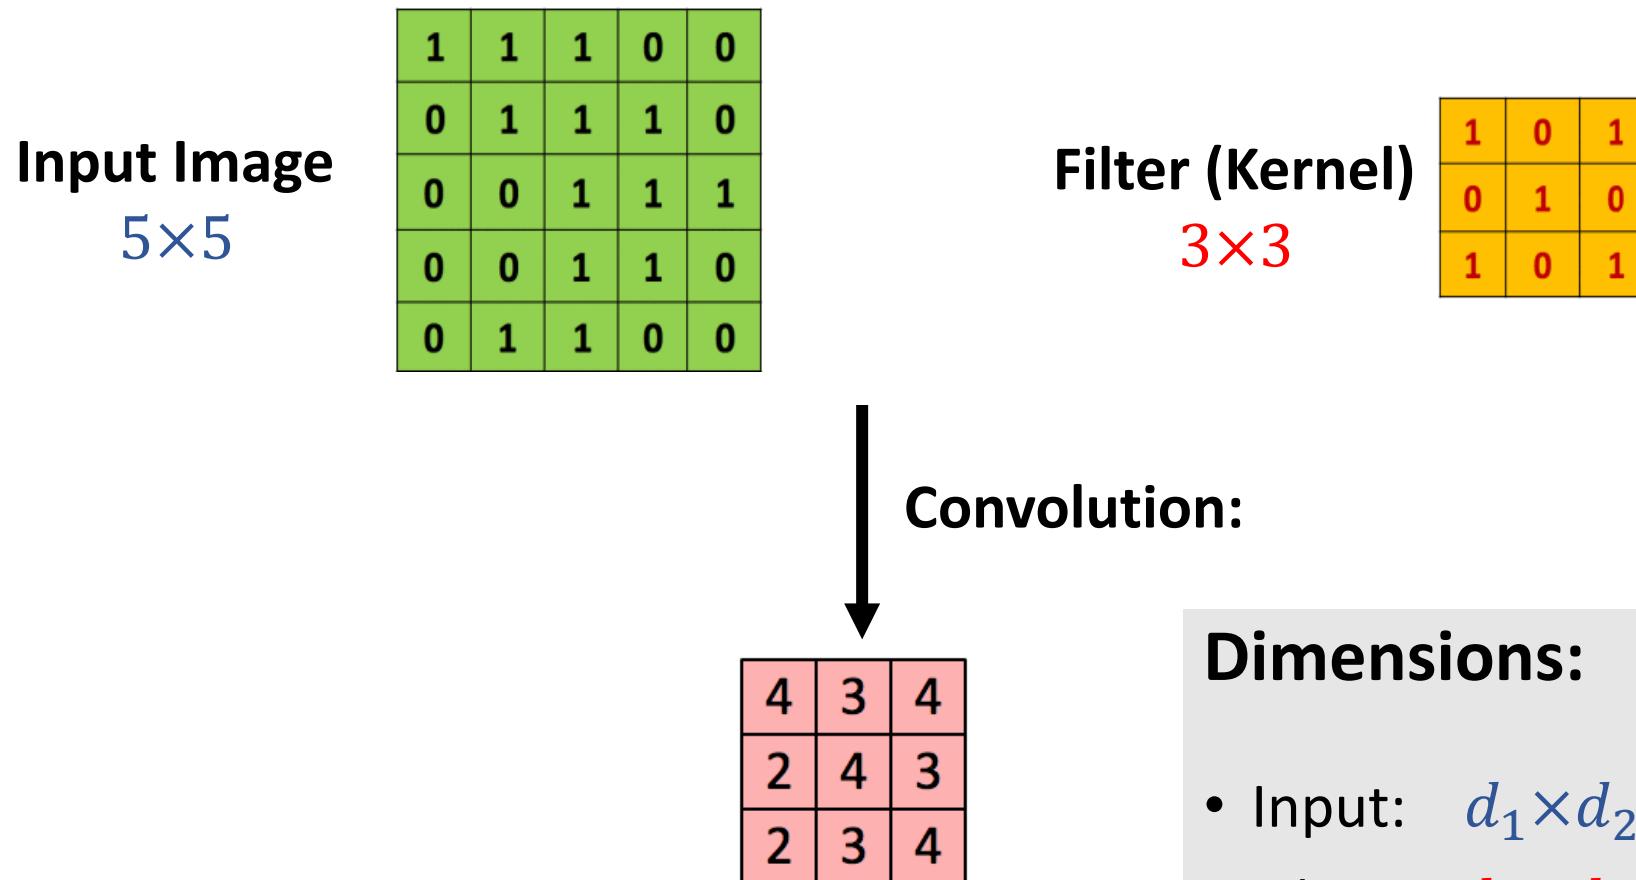


**Convolution:**

|   |   |
|---|---|
| 4 | 3 |
| 3 | 4 |

**Result**  
 $2 \times 2$

# Convolution



Dimensions:

- Input:  $d_1 \times d_2$
- Filter:  $k_1 \times k_2$
- Output:  $(d_1 - k_1 + 1) \times (d_2 - k_2 + 1)$

# Convolution

Input Image



Filters

|                |   |  |
|----------------|---|--|
| Identity<br>等式 | $\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$ | The output image is identical to the input image, showing the profile of a dog's head. This visualizes the identity filter, which preserves the original input without applying any changes. |
|----------------|---|--|

Outputs

# Convolution

Input Image



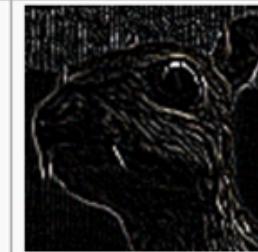
濾波器

Filters

Edge  
detection  
边缘检测

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

Outputs



- The inner product of  $\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$  and  $\begin{bmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{bmatrix}$  is **0**.
- The inner product of  $\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$  and  $\begin{bmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.2 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{bmatrix}$  is **0.8**.

# Convolution

Input Image



Filters

| Sharpen<br>锐化 | $\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$ | A small square image of the same camel's head as the input, but with more pronounced features and higher contrast, indicating it has been processed by a sharpening filter. |
|---------------|---|---|
|---------------|---|---|

Outputs

- The inner product of  $\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$  and  $\begin{bmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{bmatrix}$  is **0.1**.
- The inner product of  $\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$  and  $\begin{bmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.2 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{bmatrix}$  is **0.6**.

# Convolution

Input Image



Filters

|   |  |  |
|---|--|--|
| <b>Box blur</b><br>(normalized)         | $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$  | A blurred version of the input image, showing a smoother profile of the dog's head.  |
| <b>Gaussian blur</b><br>(approximation) | $\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$ | A slightly more blurred version of the input image compared to the box blur, showing a smoother profile of the dog's head. |

Outputs  
(or feature map)

# Convolution



Input

# Convolution: Zero Padding

Input Image  
 $5 \times 5$

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Filter (Kernel)  
 $3 \times 3$

|   |   |   |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

Convolution:

|   |            |   |            |   |            |   |   |
|---|------------|---|------------|---|------------|---|---|
| 1 | $\times_1$ | 1 | $\times_0$ | 1 | $\times_1$ | 0 | 0 |
| 0 | $\times_0$ | 1 | $\times_1$ | 1 | $\times_0$ | 1 | 0 |
| 0 | $\times_1$ | 0 | $\times_0$ | 1 | $\times_1$ | 1 | 1 |
| 0 | 0          | 1 | 1          | 1 | 0          | 0 |   |
| 0 | 1          | 1 | 0          | 0 |            |   |   |

Image

|   |  |  |
|---|--|--|
| 4 |  |  |
|   |  |  |
|   |  |  |
|   |  |  |

Convolved Feature

## Dimensions:

- Input:  $d_1 \times d_2$
- Filter:  $3 \times 3$
- Output:  $(d_1 - 2) \times (d_2 - 2)$

# Convolution: Zero Padding

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
|   |   |   |   |   |

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
|   |   |   |   |   |

**Problem:** the output is smaller than the input!

- If the input is  $28 \times 28$  and the filter is  $3 \times 3$ , you can have at most 13 convolutional layers. (The output of the 13<sup>th</sup> layer is  $2 \times 2$ .)

**Convolution:**

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

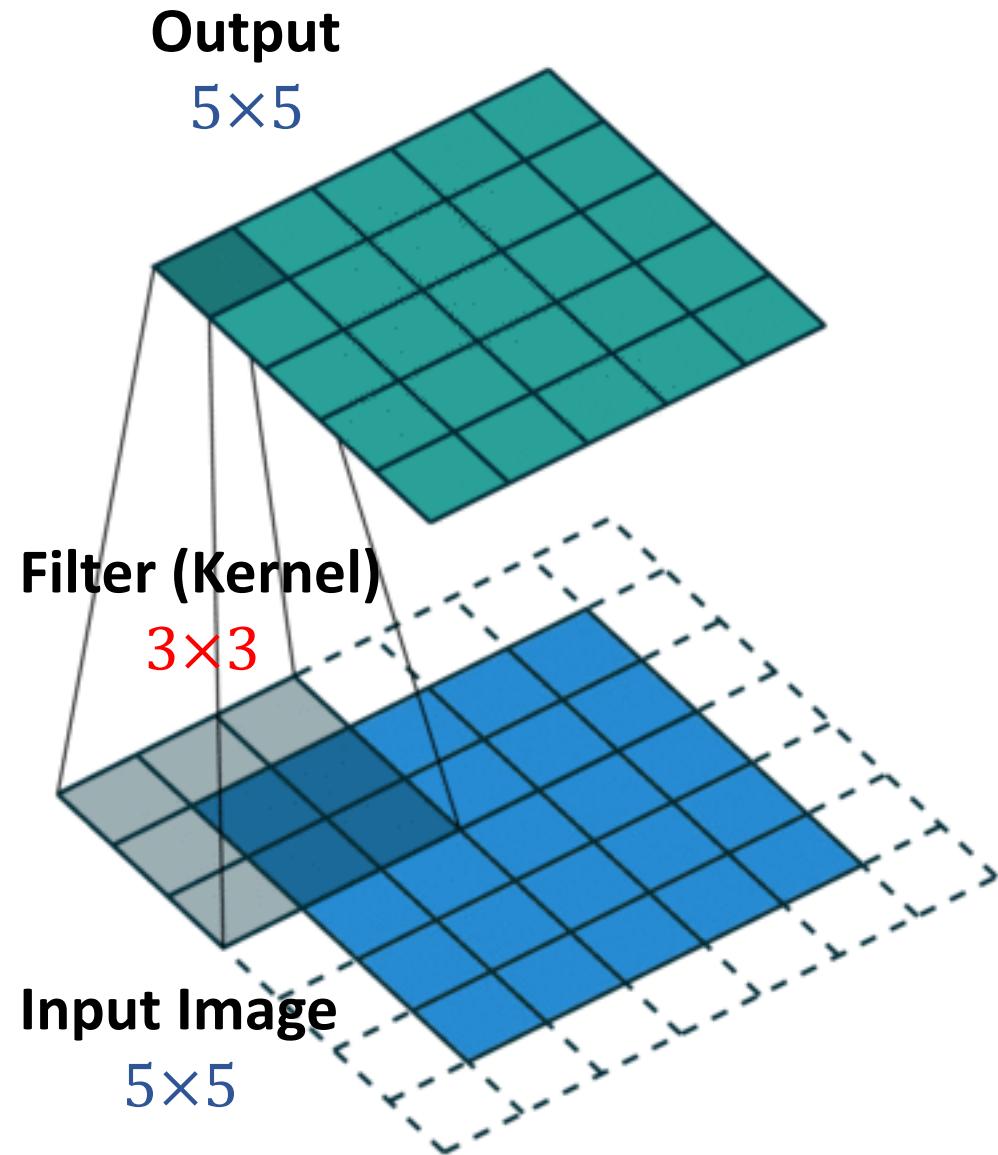
|   |  |  |
|---|--|--|
| 4 |  |  |
|   |  |  |
|   |  |  |
|   |  |  |

Convolved Feature

**Dimensions:**

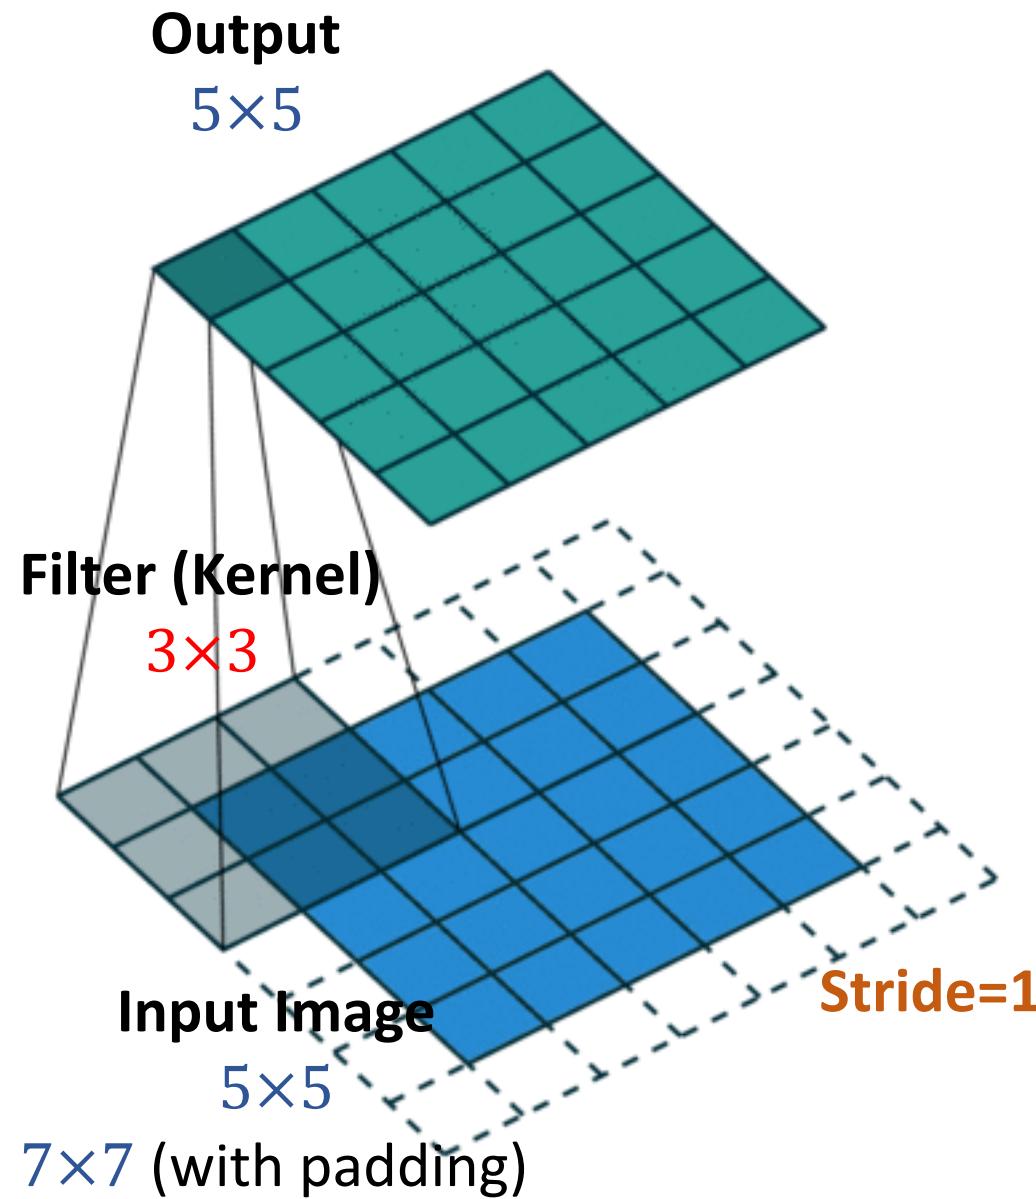
- Input:  $d_1 \times d_2$
- Filter:  $3 \times 3$
- Output:  $(d_1 - 2) \times (d_2 - 2)$

# Convolution: Zero Padding



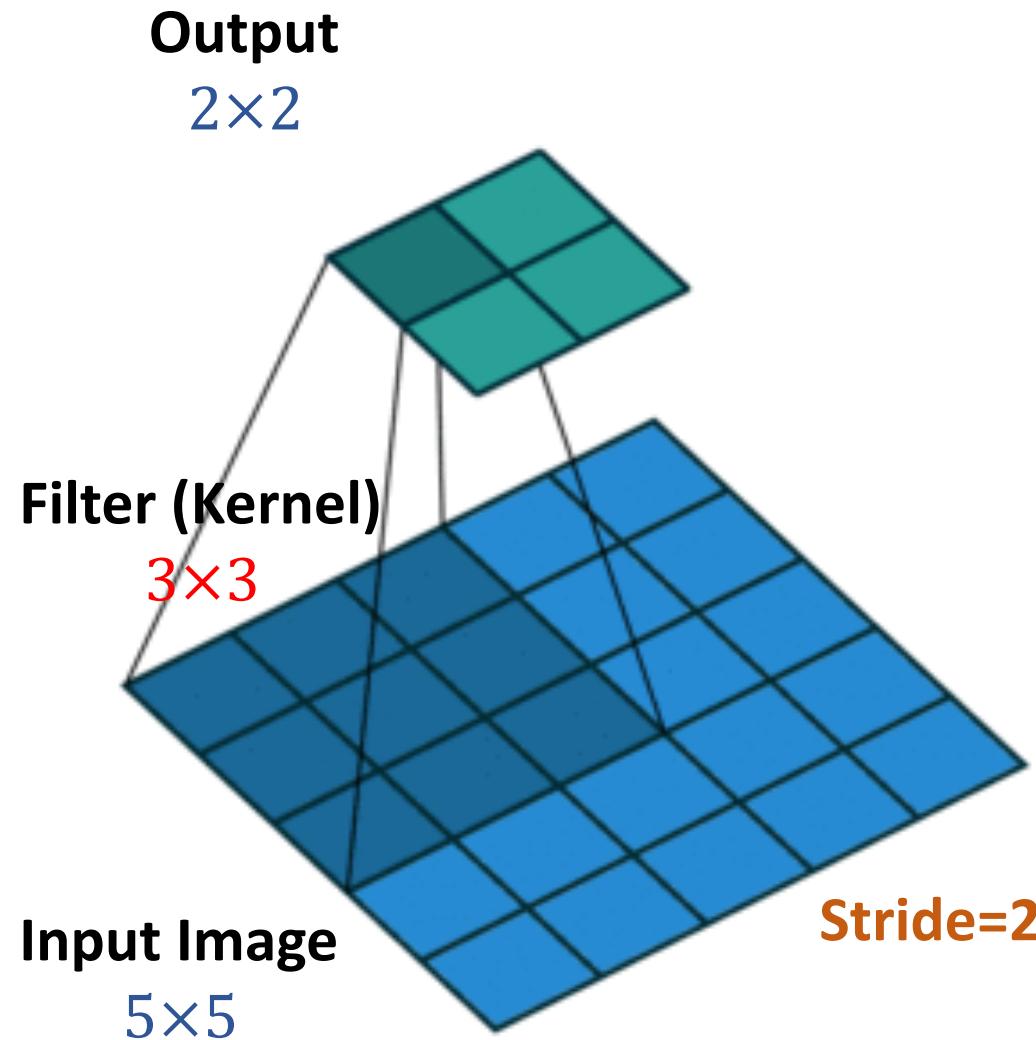
- Add a “boarder” of all-zeros.
- Increase the input shape:
  - From  $d_1 \times d_2$  to  $(d_1 + 2) \times (d_2 + 2)$ .
  - If the filter is  $3 \times 3$ , the output is  $d_1 \times d_2$ .

# Convolution: Stride



- In the previous examples, the **stride** is **1**.
  - The filter moves **1** step each time.

# Convolution: Stride

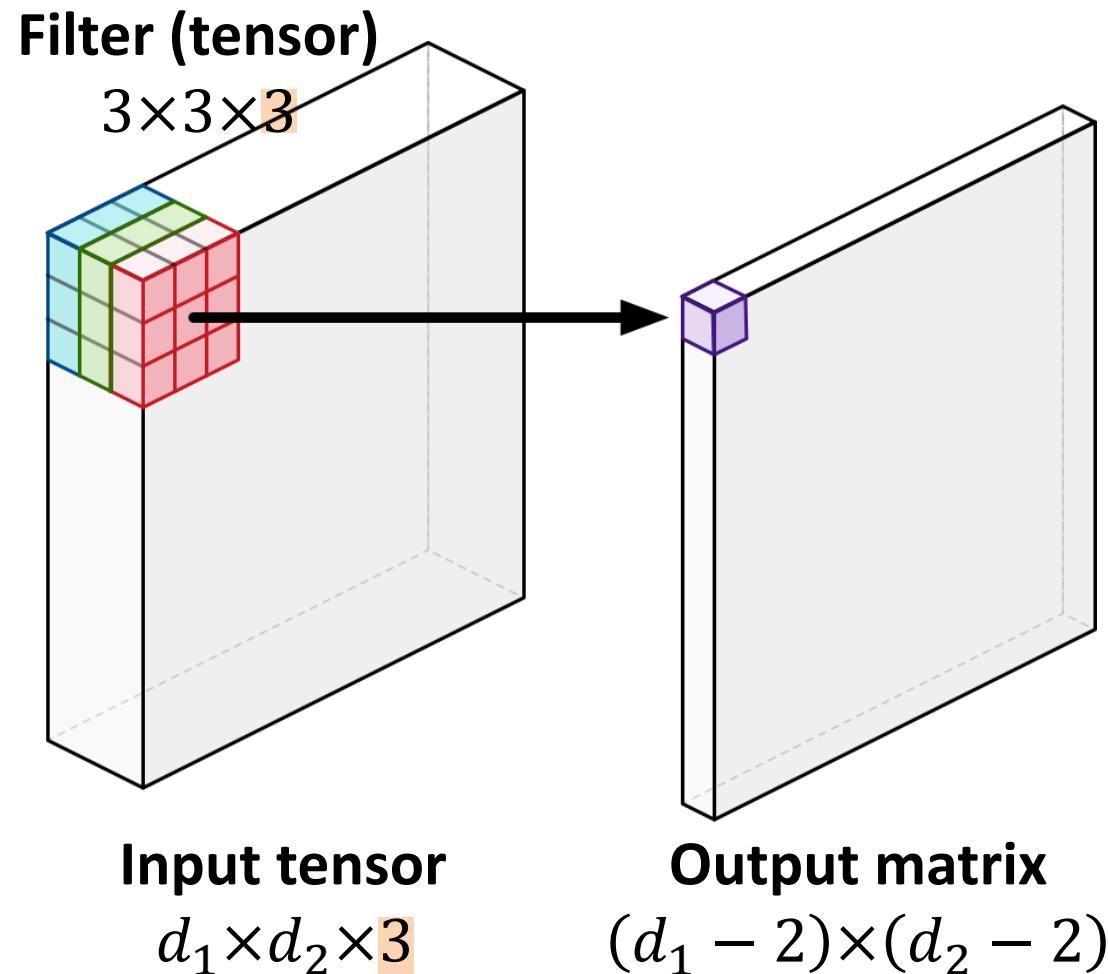


- In the previous examples, the **stride** is **1**.
  - The filter moves **1** step each time.
- The **stride** can be **2** or even larger.

## Dimensions:

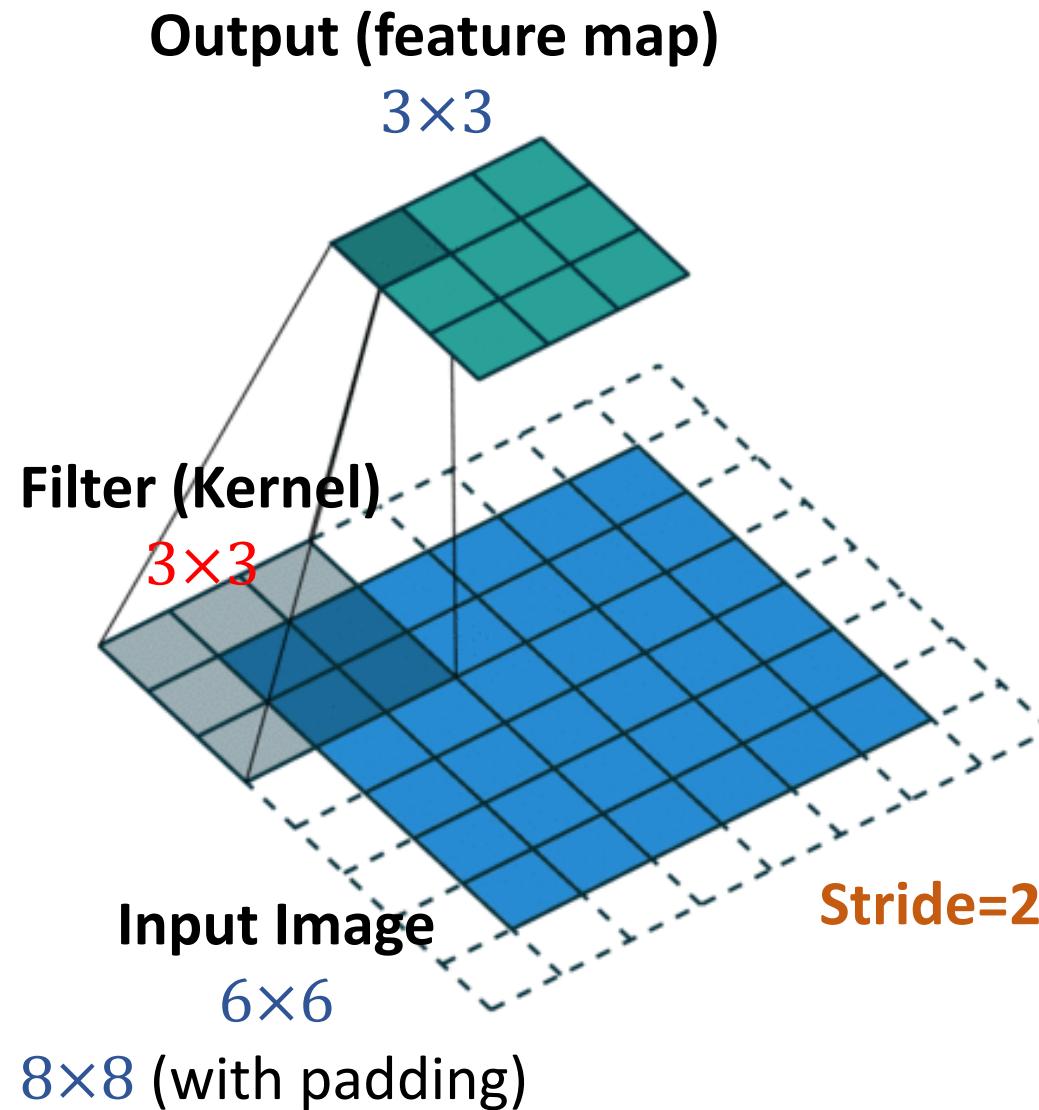
- Input:  $d_1 \times d_2$
- Filter:  $k_1 \times k_2$
- Stride:  $s$
- Output:  $\left(\left\lfloor \frac{d_1 - k_1}{s} \right\rfloor + 1\right) \times \left(\left\lfloor \frac{d_2 - k_2}{s} \right\rfloor + 1\right)$

# Convolution for Tensors



- The input is the order-3 tensor of shape  $d_1 \times d_2 \times 3$ , e.g., a color image.
- The filter is a order-3 tensor of shape  $3 \times 3 \times 3$ .
- The output is a matrix (order-2 tensor) of shape  $(d_1 - 2) \times (d_2 - 2)$ .

# Convolution: Key Concepts

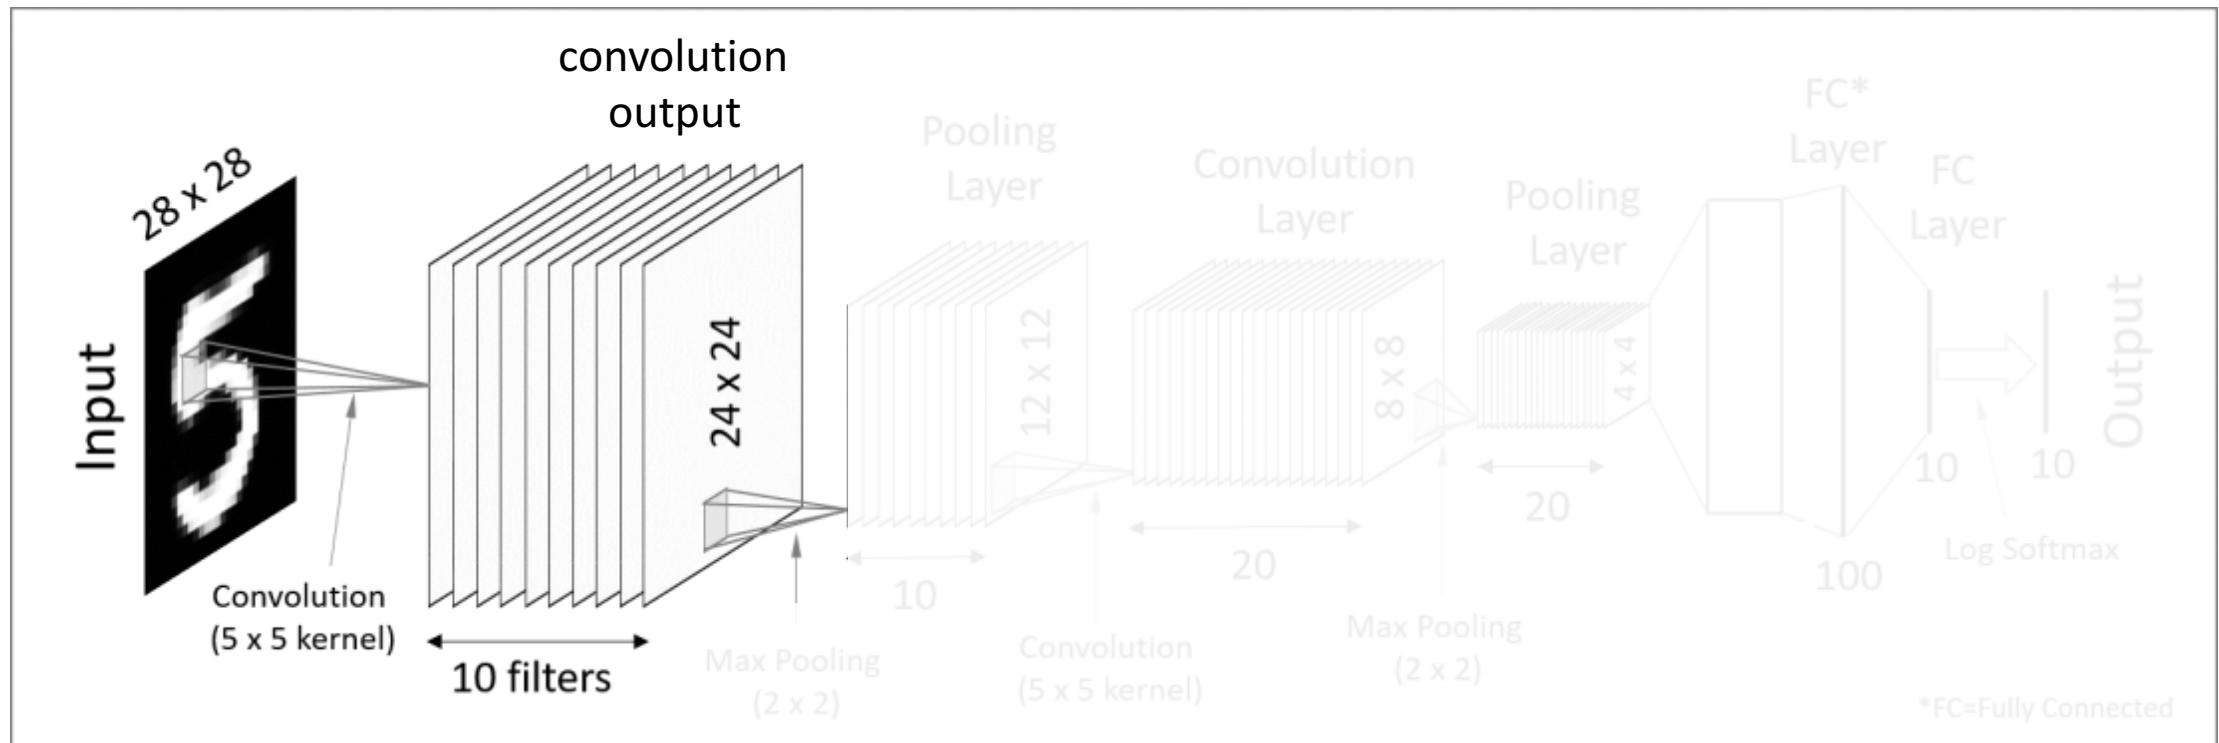


- Input, filter (kernel), and output (feature map).
- Convolution for matrix (order-2 tensor).
- Convolution for order-3 tensor.
- Zero padding.
- Stride.

# **Convolutional Neural Networks (CNNs)**

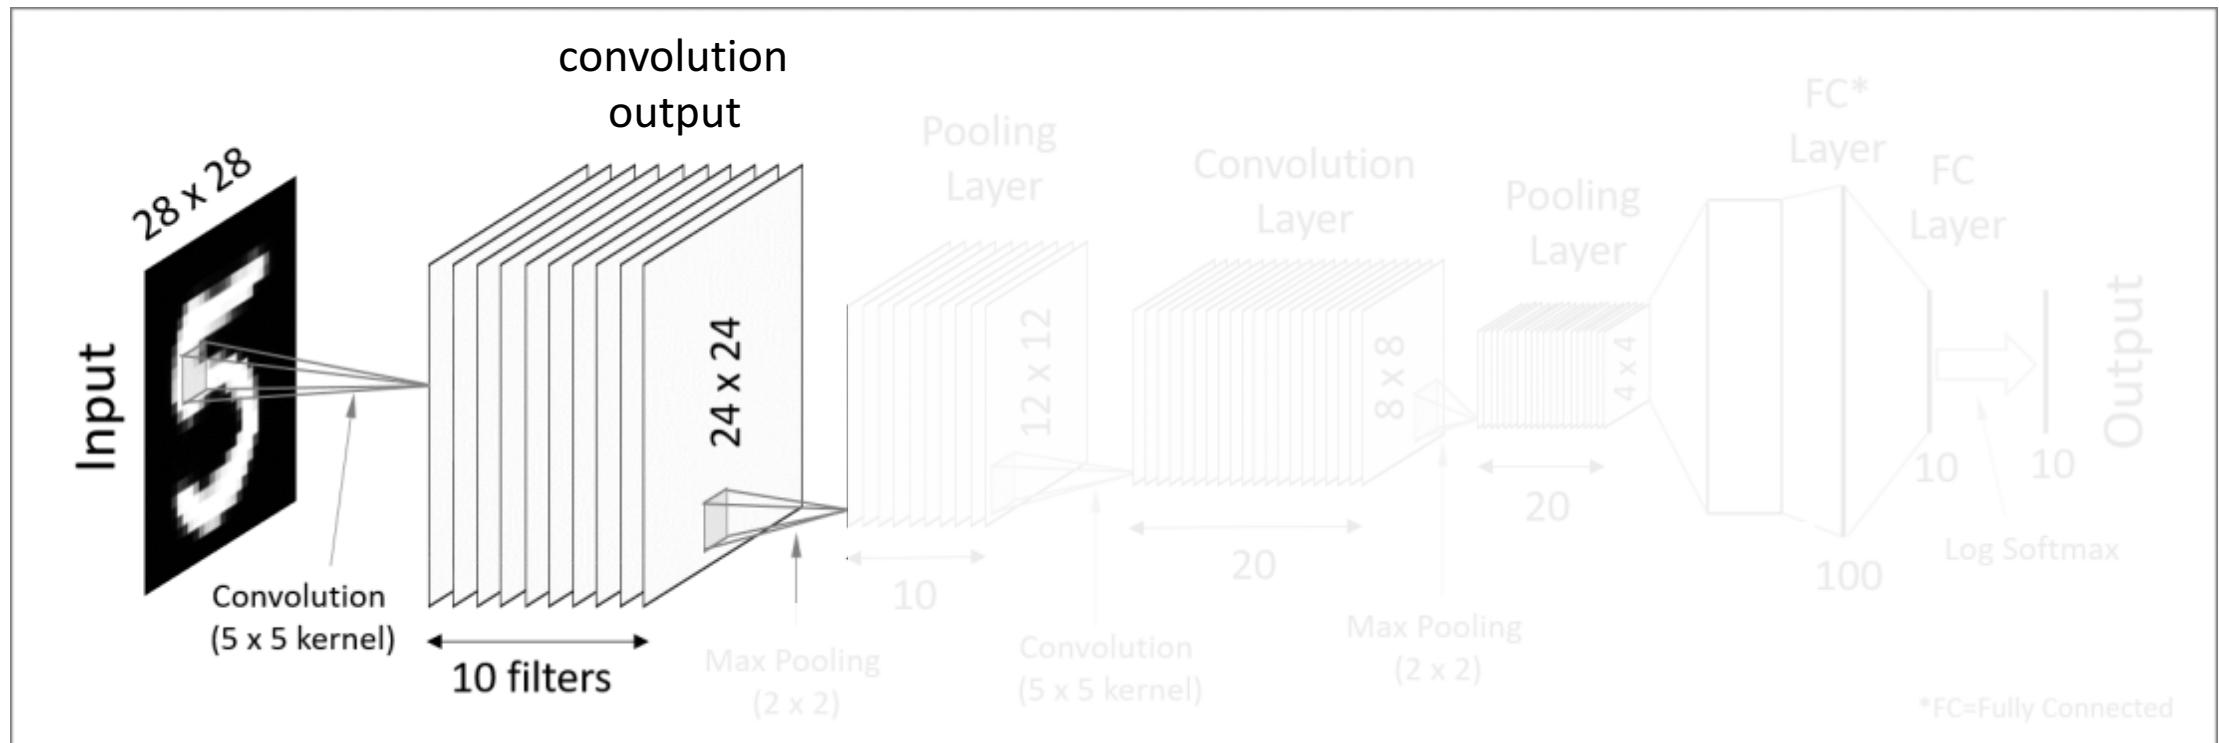
# CNN: Convolutional Layer 1

- Use multiple filters (kernels) for the feature mapping. E.g.,
  - using ten  $5 \times 5$  filters, (merely  $10 \times 5 \times 5 = 250$  parameters!)
  - the results are ten  $24 \times 24$  matrices.



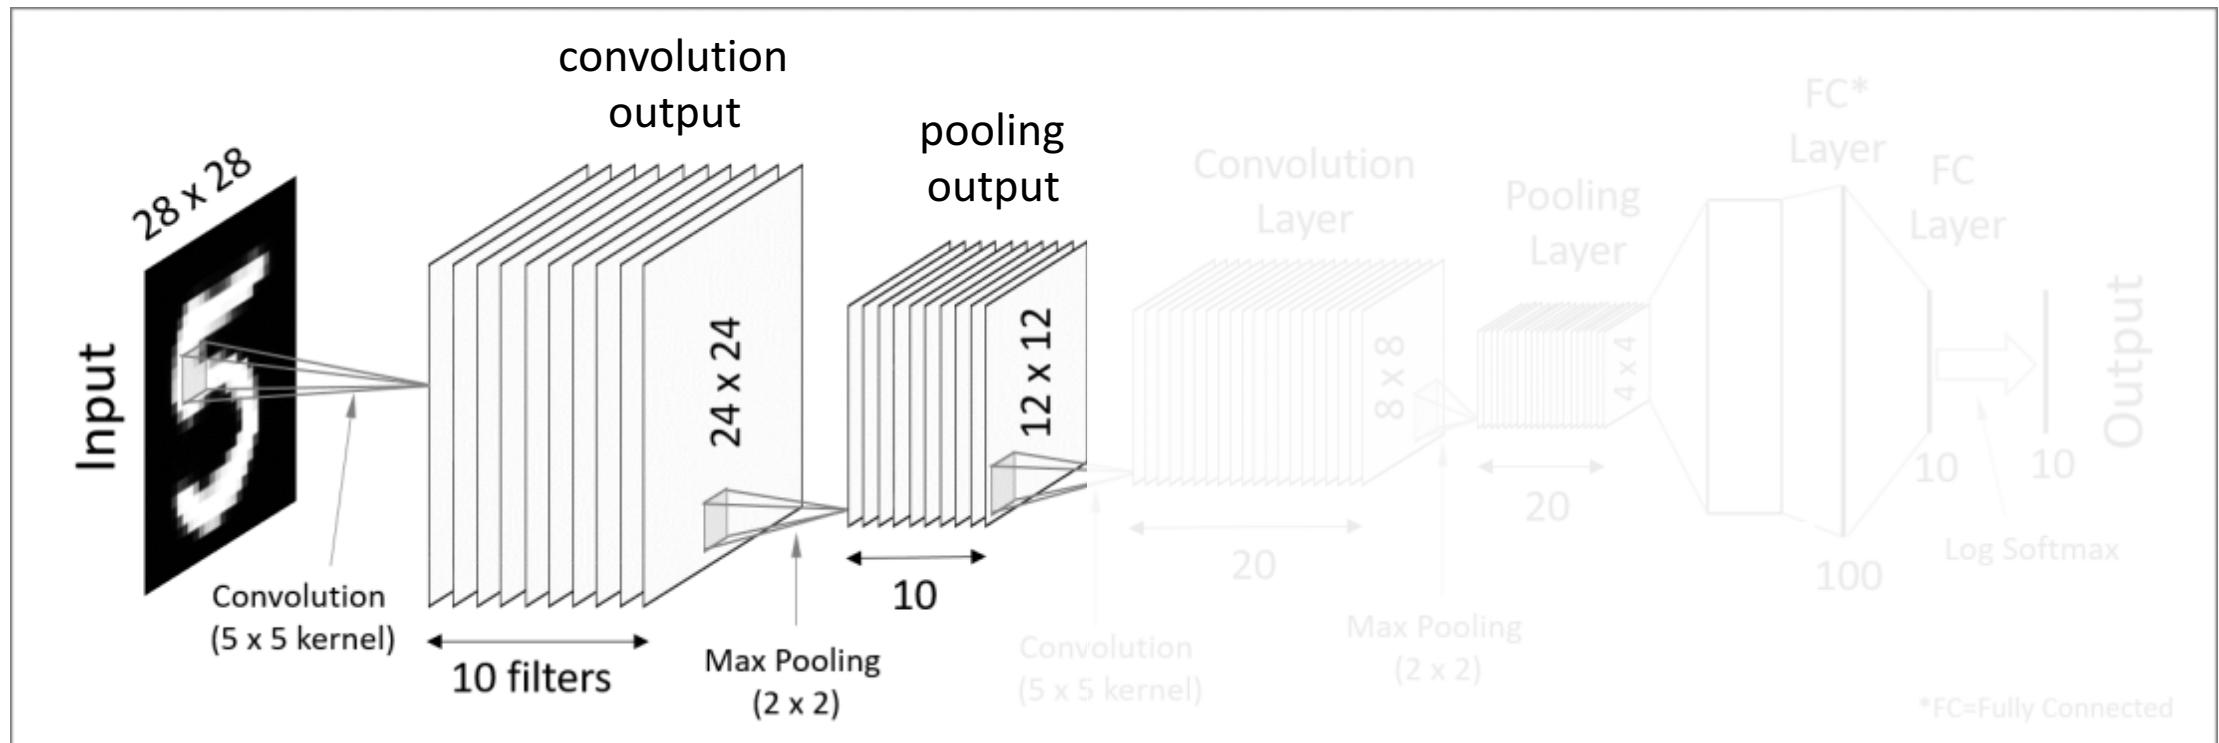
# CNN: Convolutional Layer 1

- Use multiple filters (kernels) for the feature mapping.
- Apply an activation function (e.g., ReLU).
- A Convolutional layer = convolution + activation.



# CNN: Pooling Layer 1

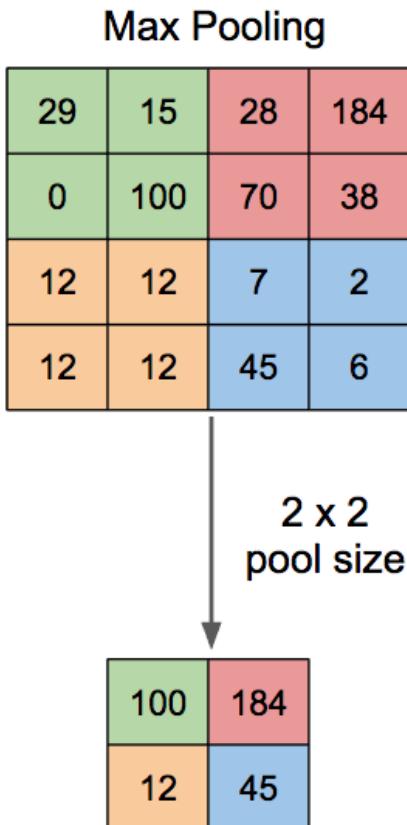
- Pooling reduces the dimensionality of each feature map.
- E.g., Max Pooling, Average Pooling, etc.



# Pooling

池化降低了每个特征映射的维数

- Pooling reduces the dimensionality of each feature map.
- E.g., Max Pooling, Average Pooling, etc.

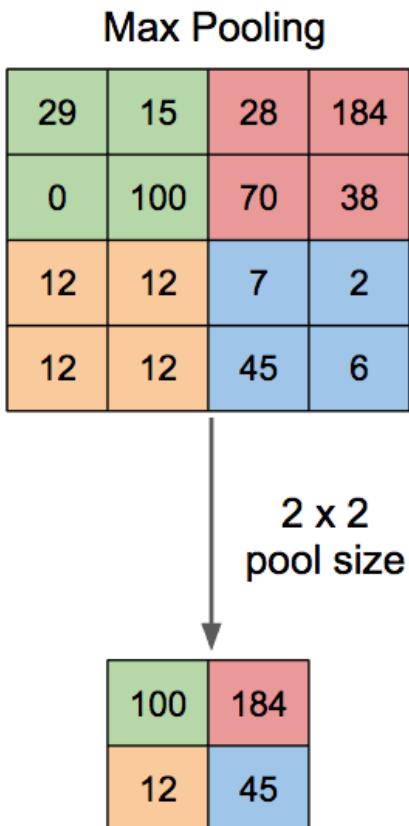


## Settings:

- Max Pooling.
- Pool size =  $2 \times 2$ .
- No overlap (stride =  $2 \times 2$ ).

# Pooling

- Pooling reduces the dimensionality of each feature map.
- E.g., Max Pooling, Average Pooling, etc.



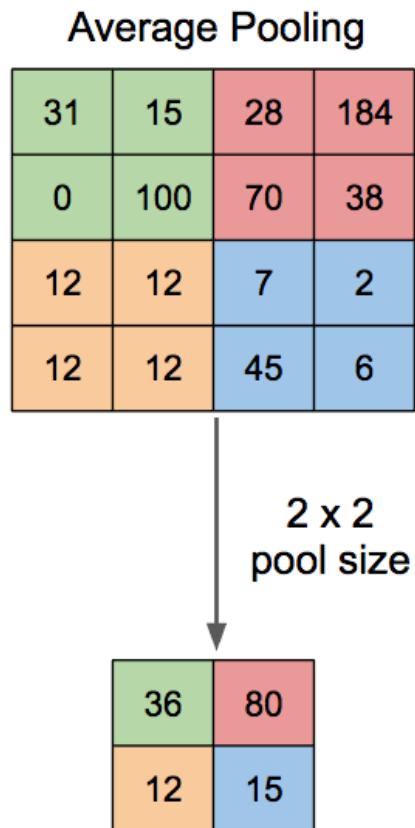
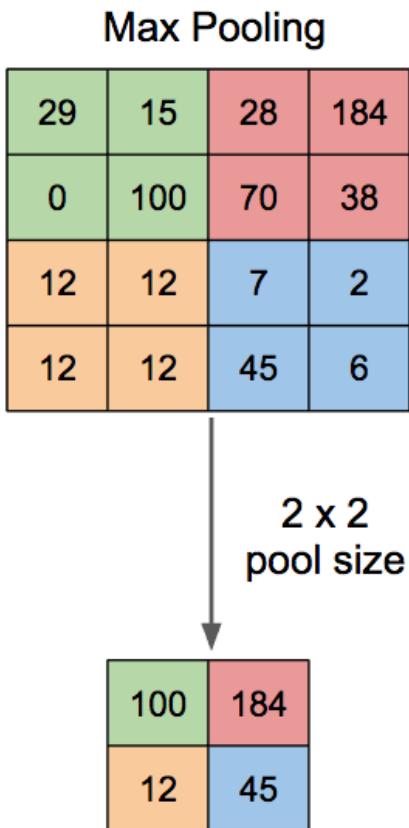
## Settings:

- Max Pooling.
- Pool size =  $2 \times 2$ .
- No overlap (stride =  $2 \times 2$ ).

**Question:** What if the stride is  $1 \times 1$ ?

# Pooling

- Pooling reduces the dimensionality of each feature map.
- E.g., Max Pooling, Average Pooling, etc.



## Settings:

- Average Pooling.
- Pool size =  $2 \times 2$ .
- No overlap (stride =  $2 \times 2$ ).

# Pooling

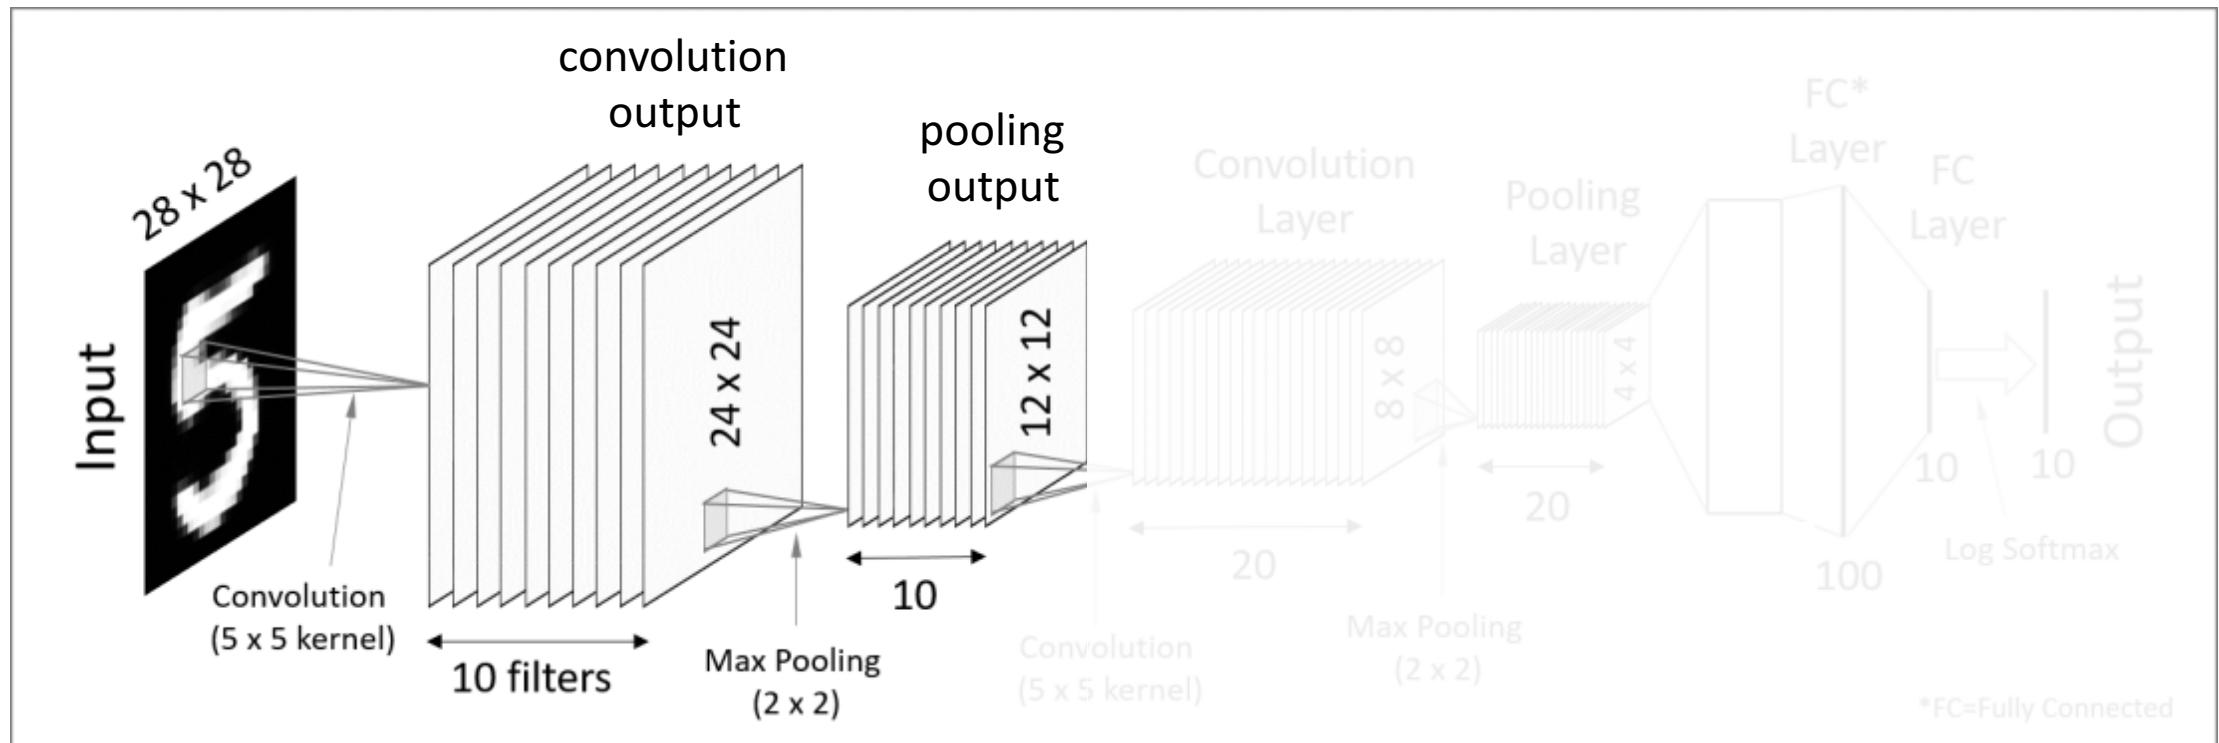
- Pooling reduces the dimensionality of each feature map.
- E.g., Max Pooling, Average Pooling, etc.
- **Question:** What is the output shape?
  - Input shape:  $12 \times 12$ .
  - Pool size:  $3 \times 3$ .
  - No overlap (stride:  $3 \times 3$ ).

# Pooling

- Pooling reduces the dimensionality of each feature map.
- E.g., Max Pooling, Average Pooling, etc.
- **Question:** What is the output shape?
  - Input shape:  $12 \times 12$ .
  - Pool size:  $3 \times 3$ .
  - With overlap (stride:  $1 \times 1$ ).

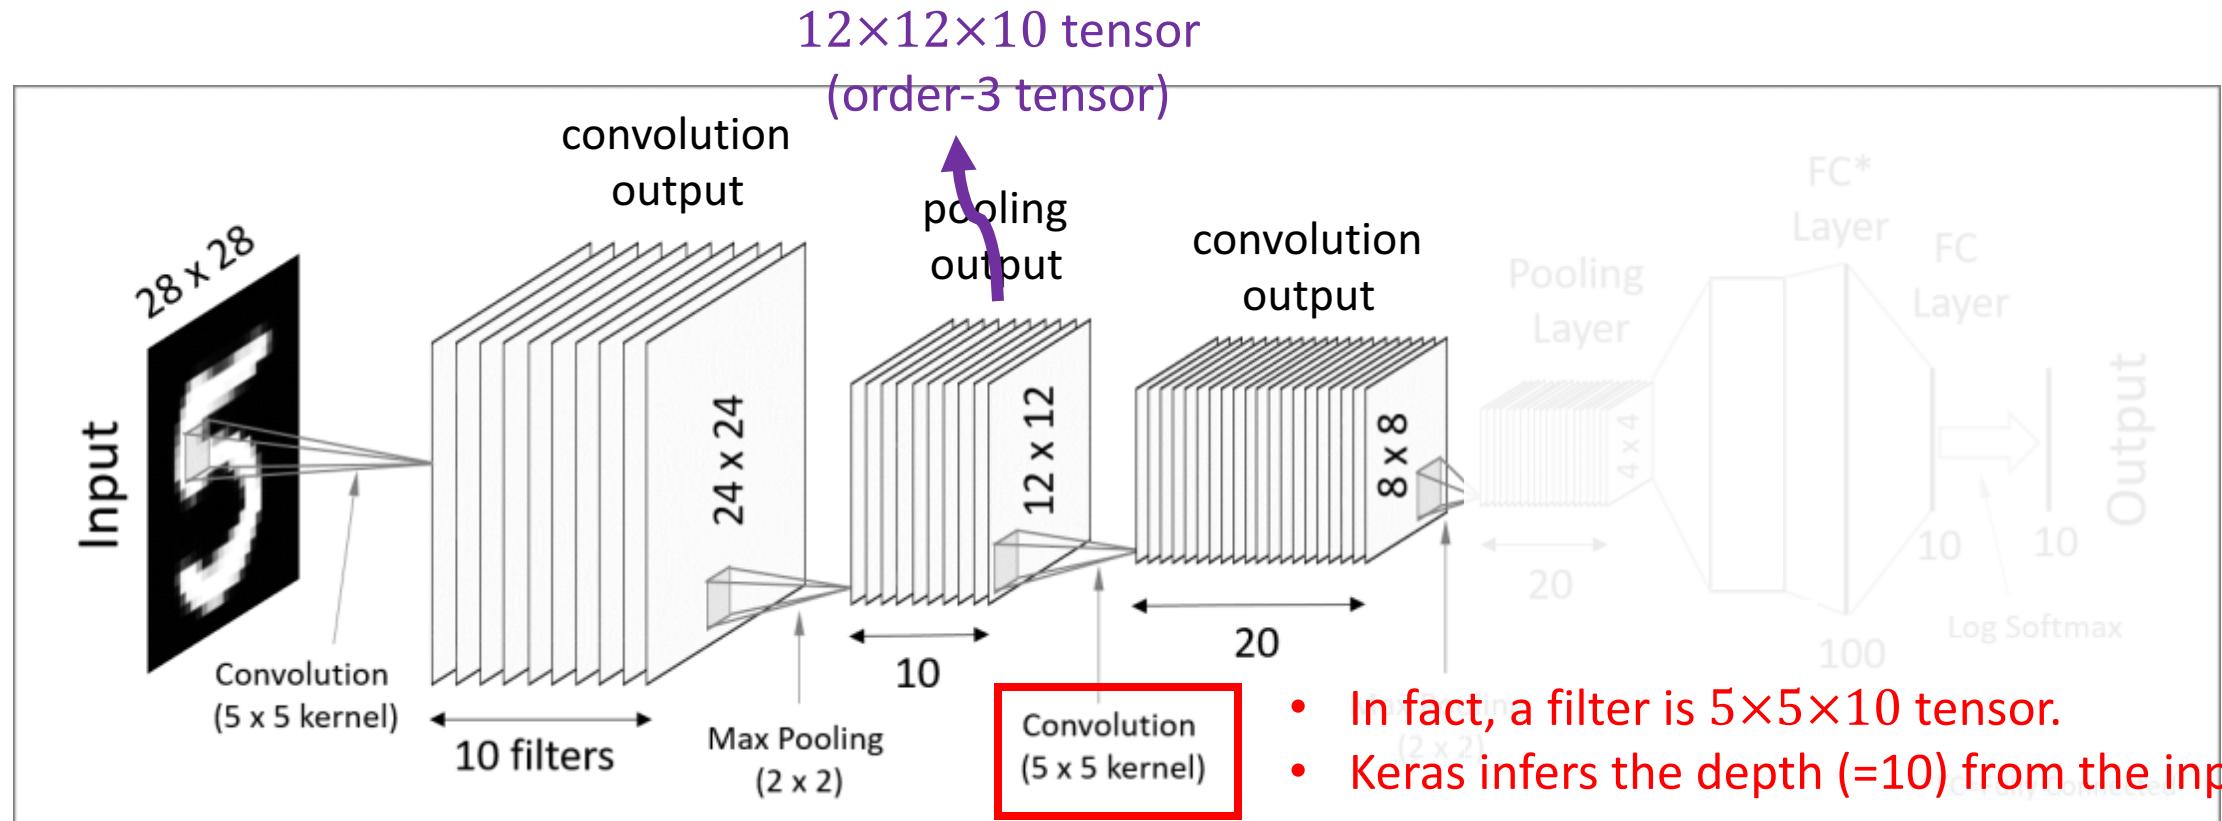
# CNN: Pooling Layer 1

- Pooling reduces the dimensionality of each feature map.
- E.g., Max Pooling, Average Pooling, etc.

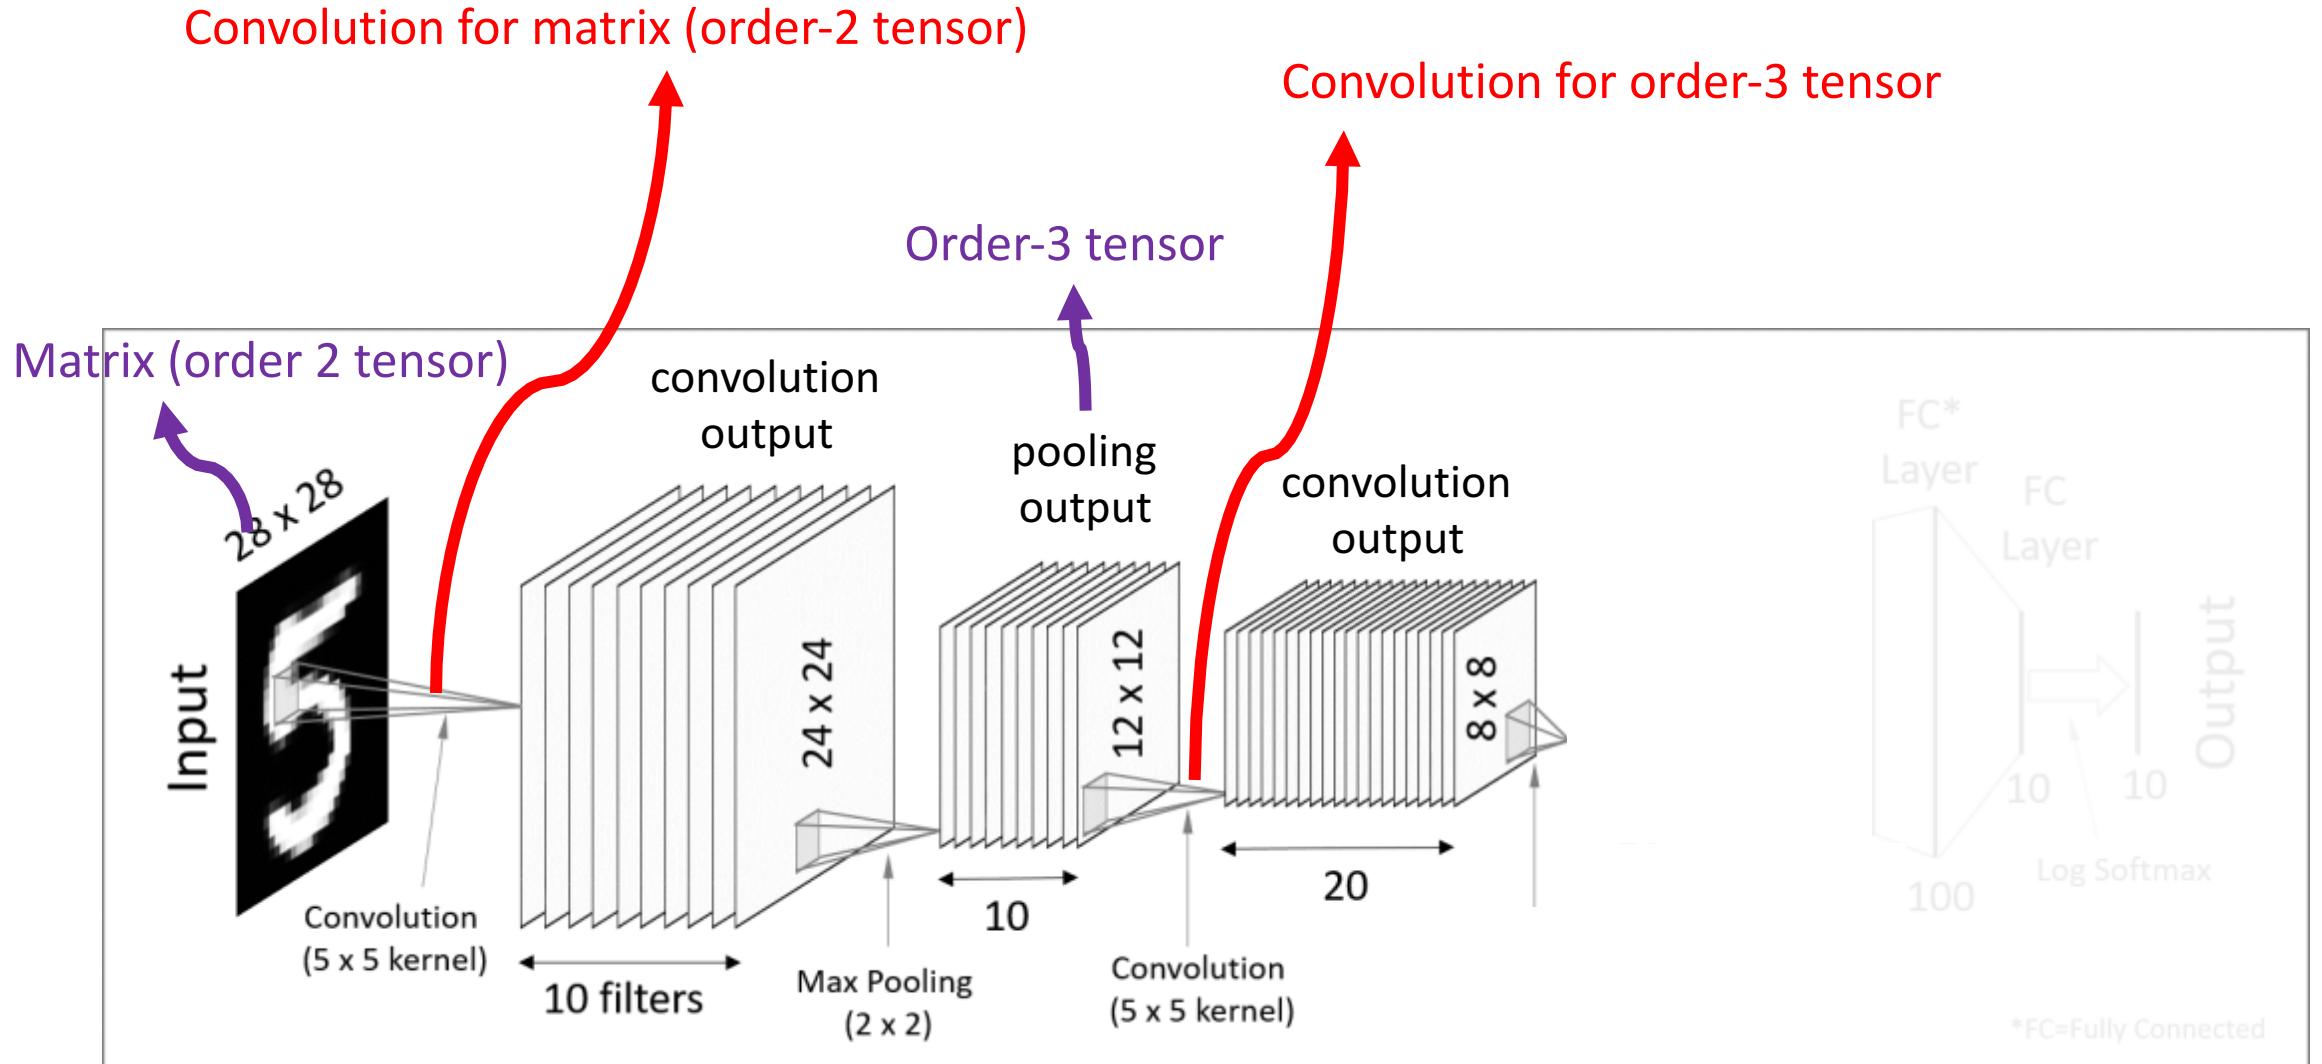


# CNN: Convolutional Layer 2

- Perform tensor convolution.
- Apply  $5 \times 5 \times 10$  tensor filter to the  $12 \times 12 \times 10$  tensor input.
- Twenty such filters → twenty  $8 \times 8$  feature maps.

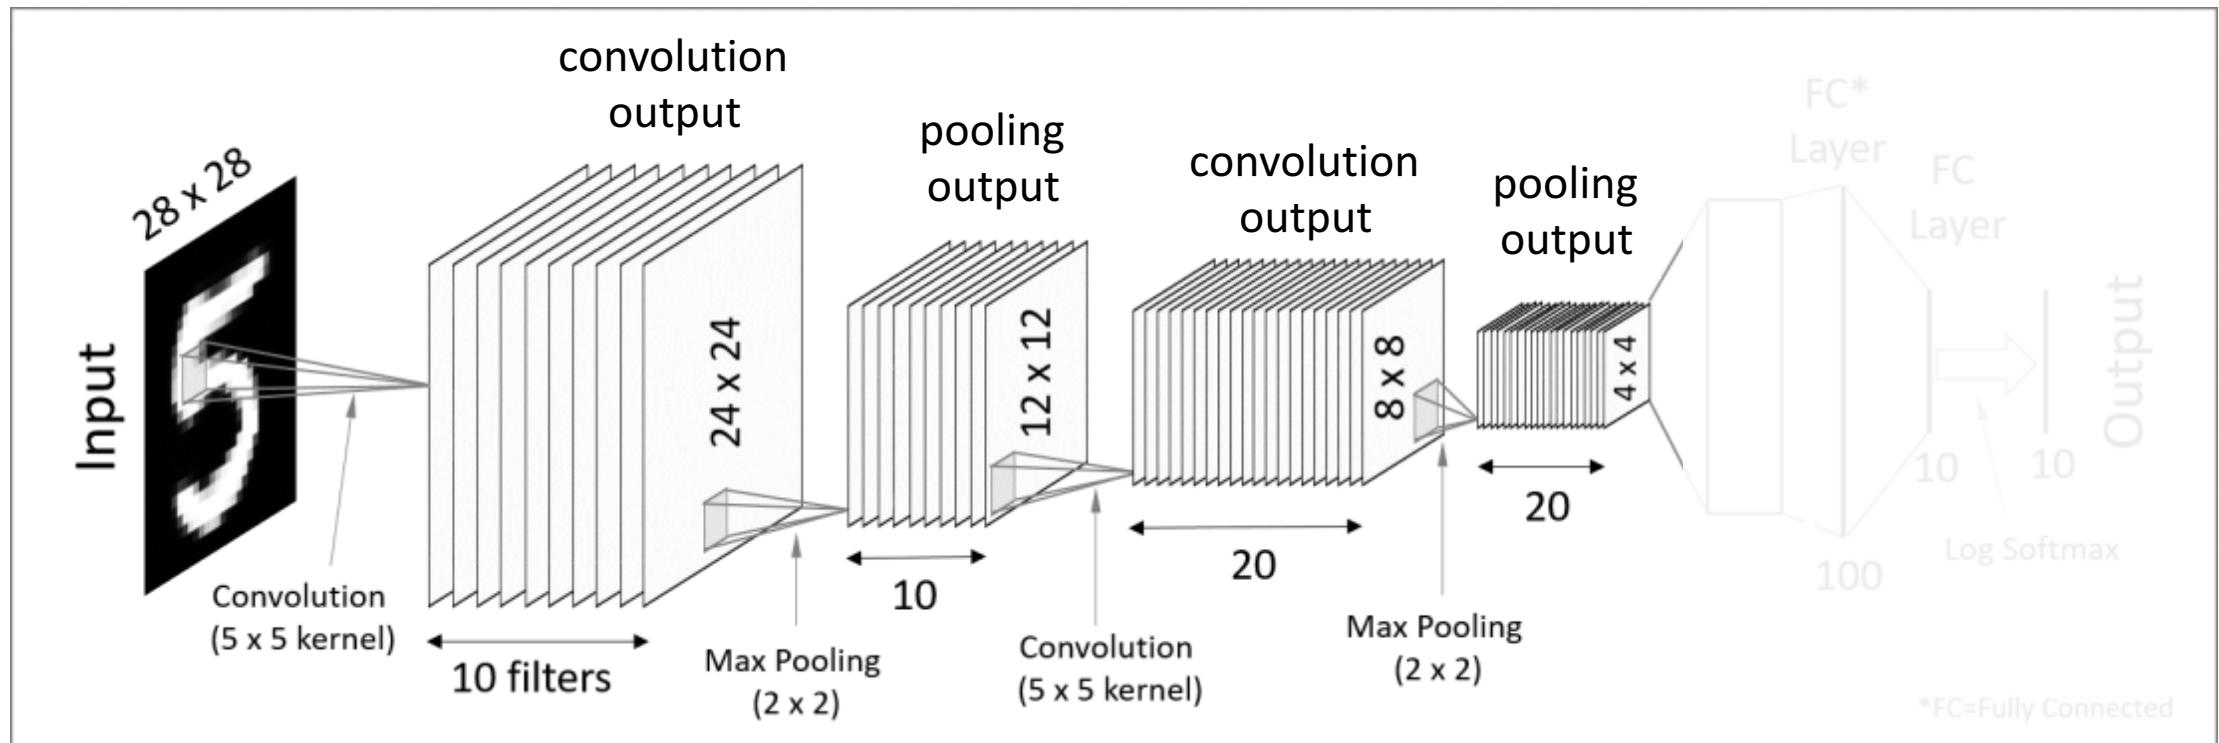


# CNN: Convolutional Layers



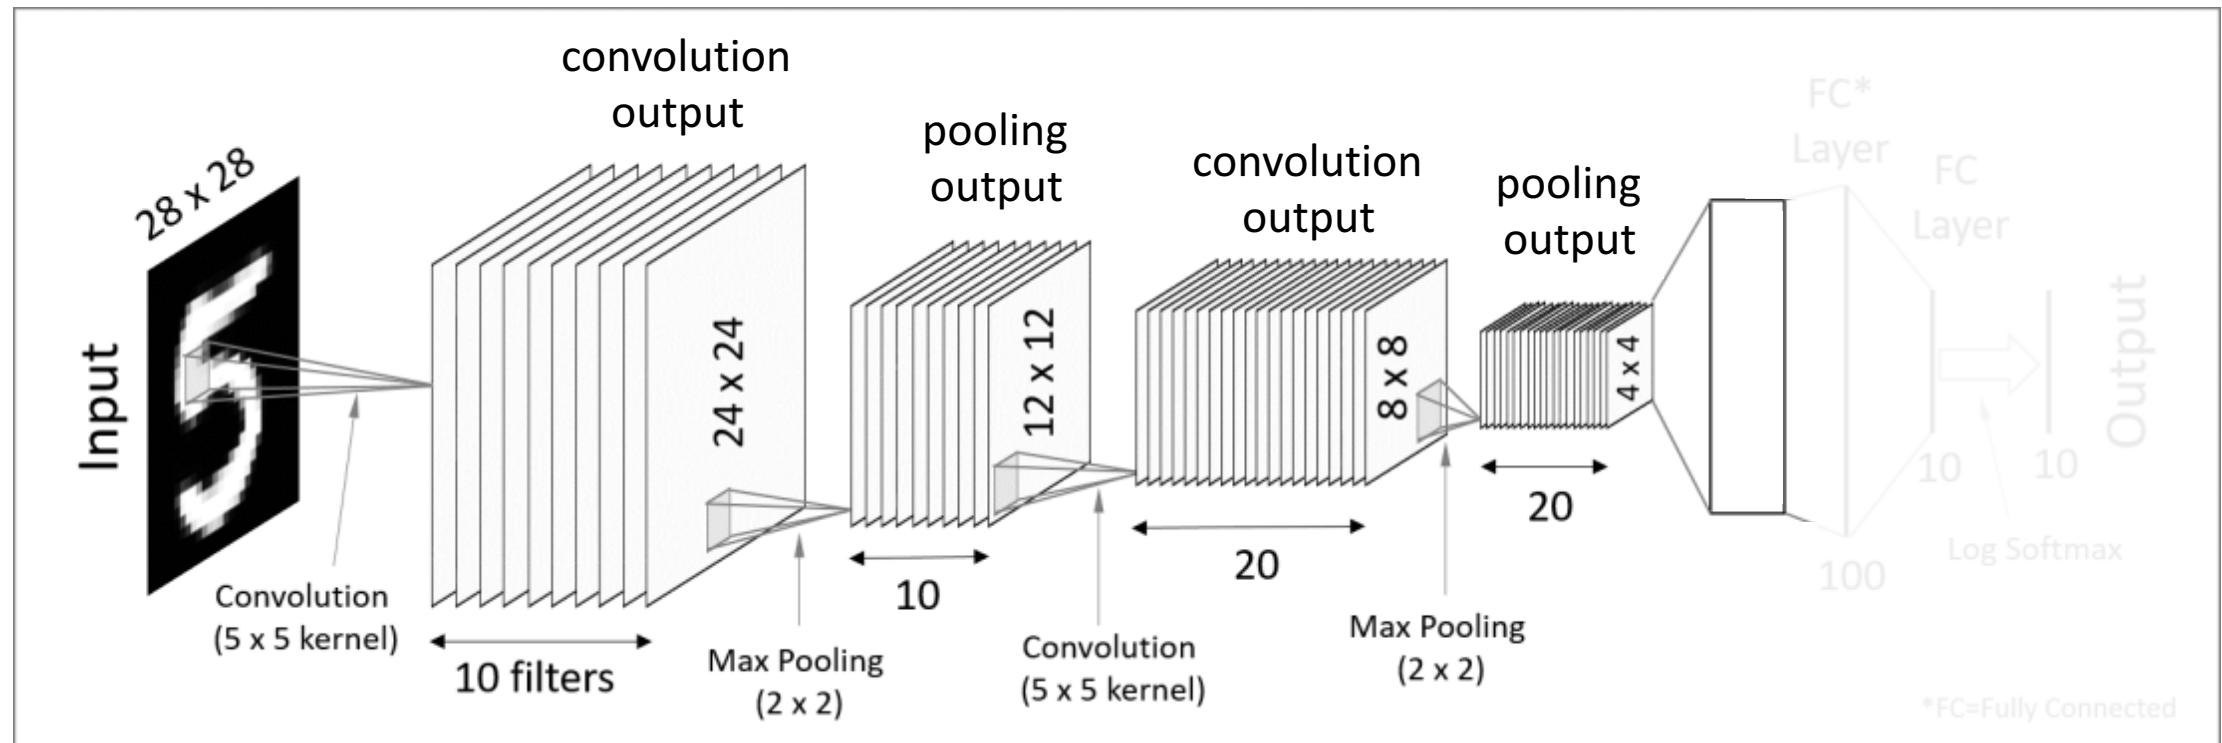
# CNN: Pooling Layer 2

- $2 \times 2$  Max Pooling applied to an  $8 \times 8$  matrix  $\rightarrow$  a  $4 \times 4$  matrix.



# CNN: Flatten (Vectorization)

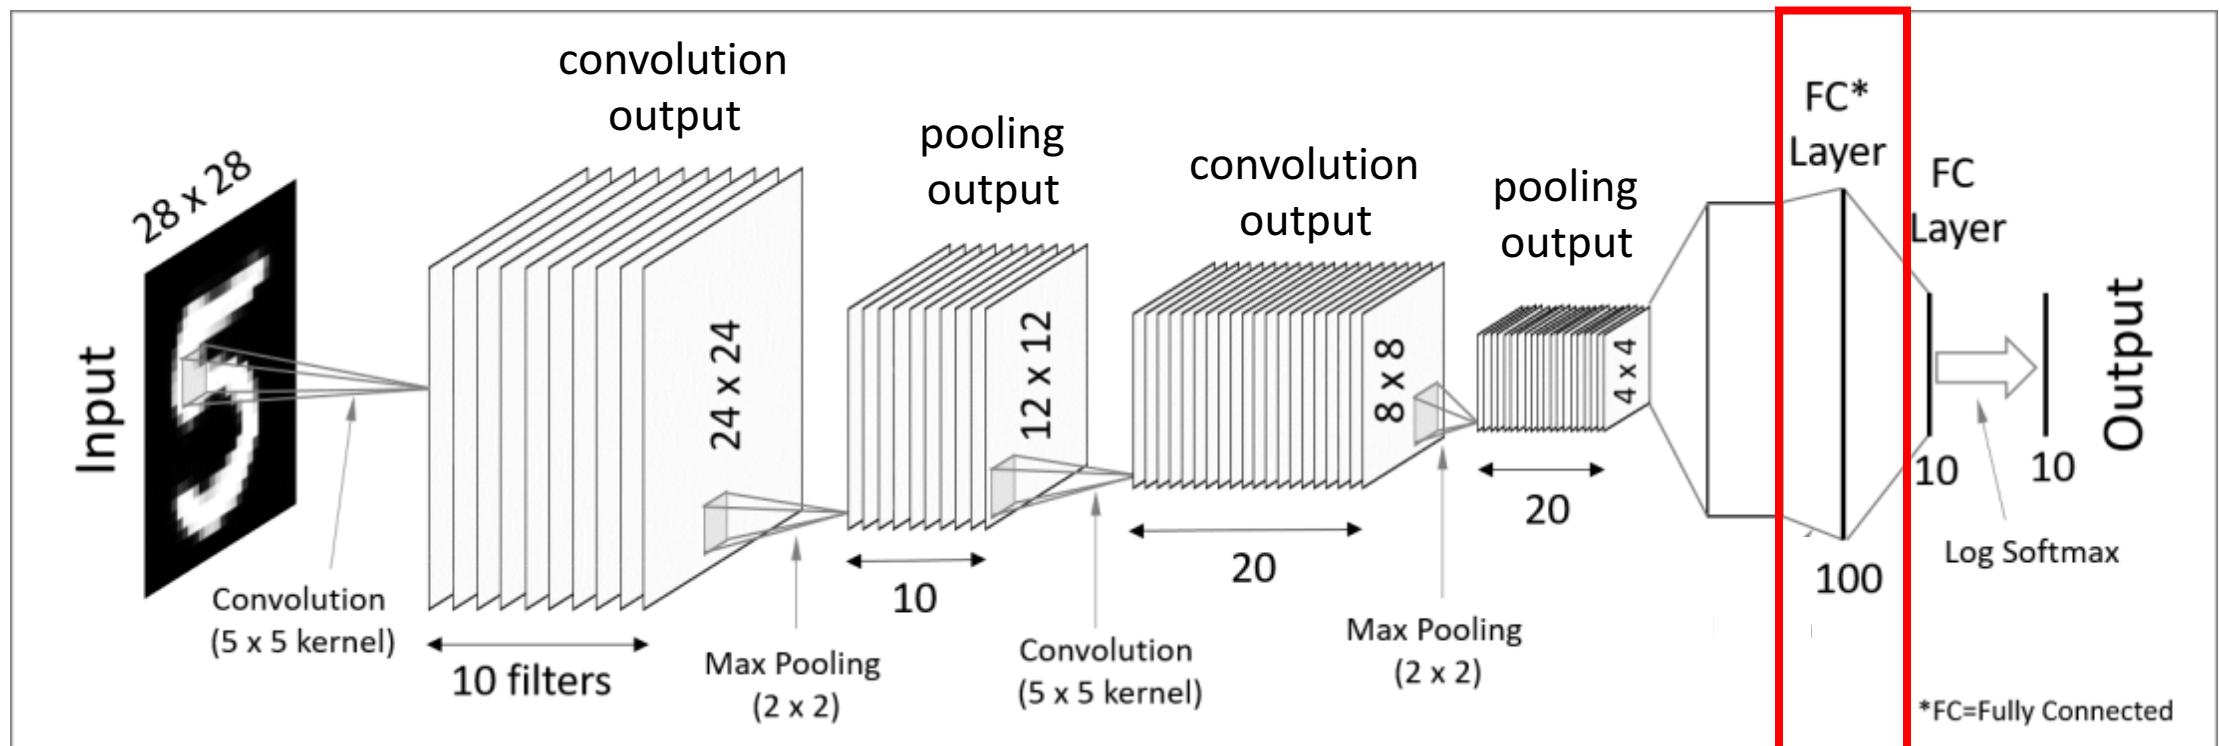
- Flatten the  $4 \times 4 \times 20$  tensor.
- Get a 320-dim vector.



# CNN: Fully-Connected Layers

- Add one or more fully-connected layers

This one is optional.



# CNN: How Many Trainable Parameters?

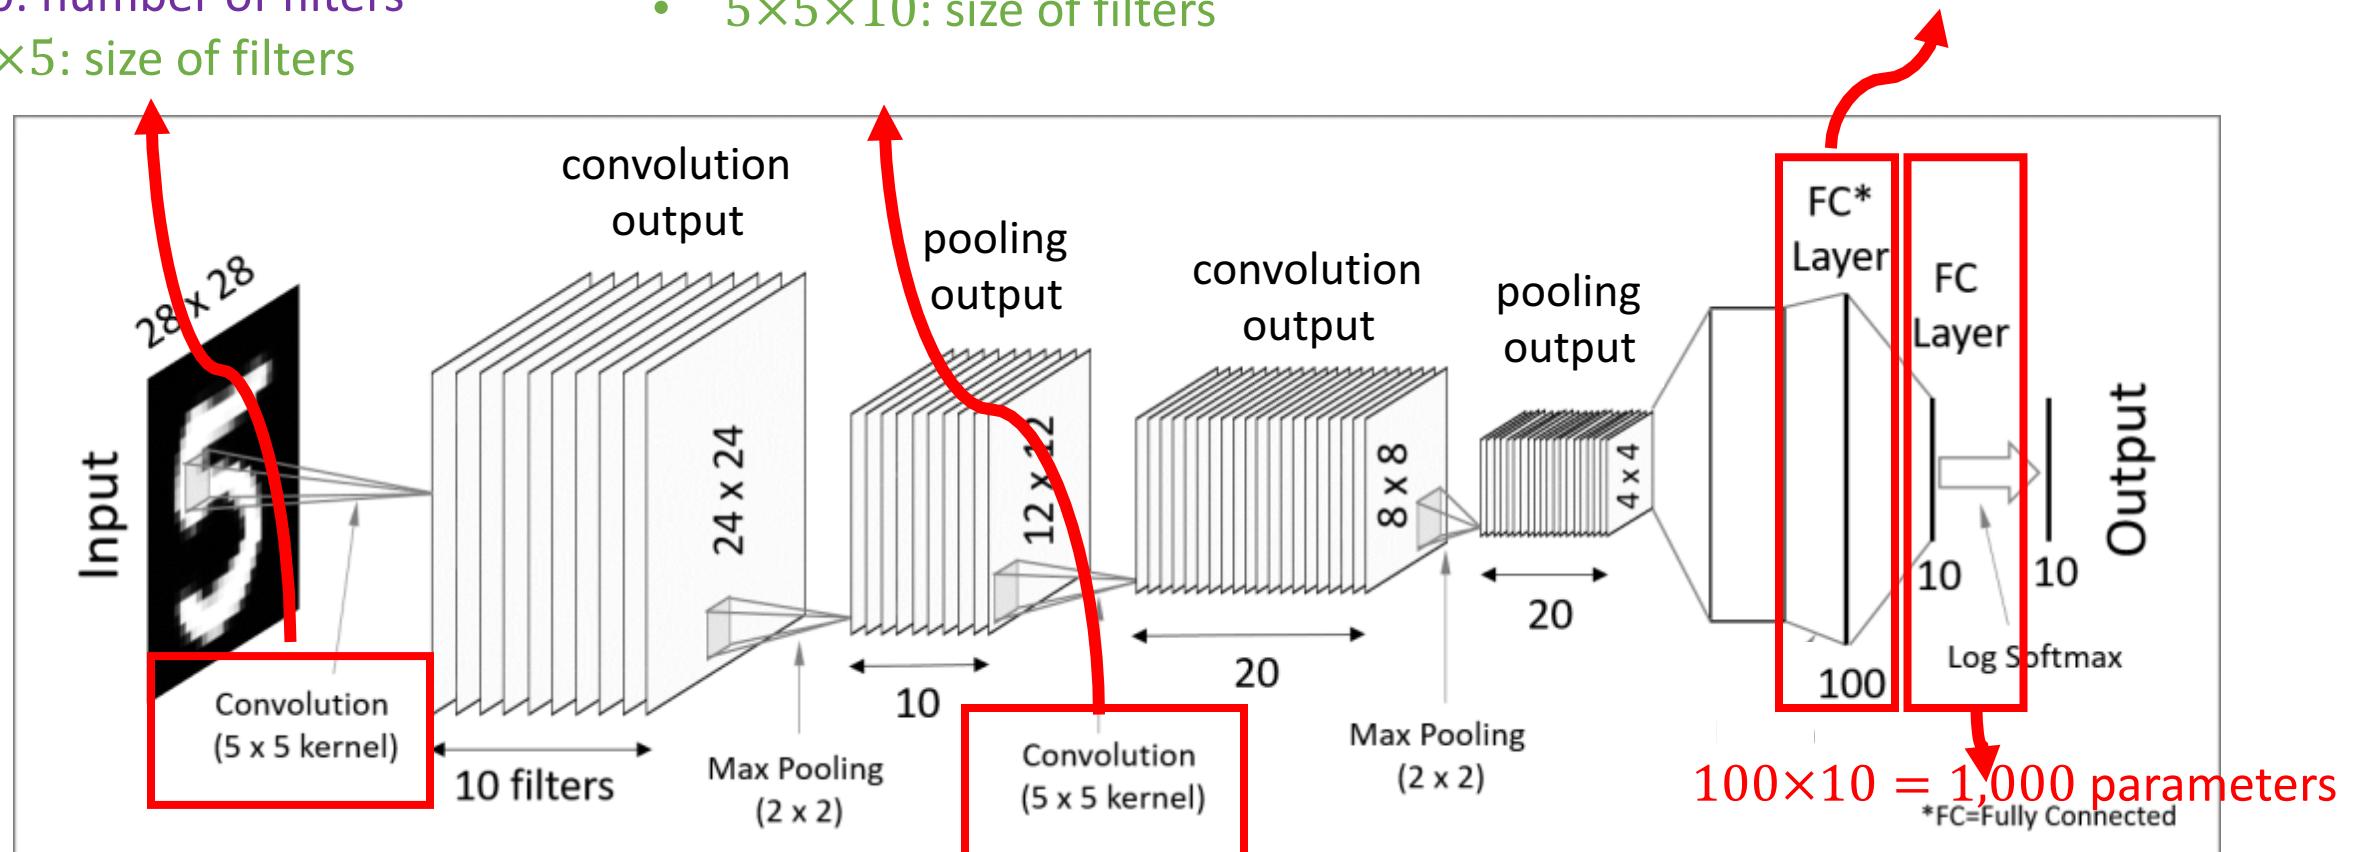
$$10 \times 5 \times 5 = 250 \text{ parameters}$$

- 10: number of filters
- 5×5: size of filters

$$20 \times 5 \times 5 \times 10 = 5000 \text{ parameters}$$

- 20: number of filters
- 5×5: size of filters

$$320 \times 100 = 32,000 \text{ parameters}$$



# CNN: How Many Trainable Parameters?

Answer:  $> 38,250$  parameters (not counting the intercepts).

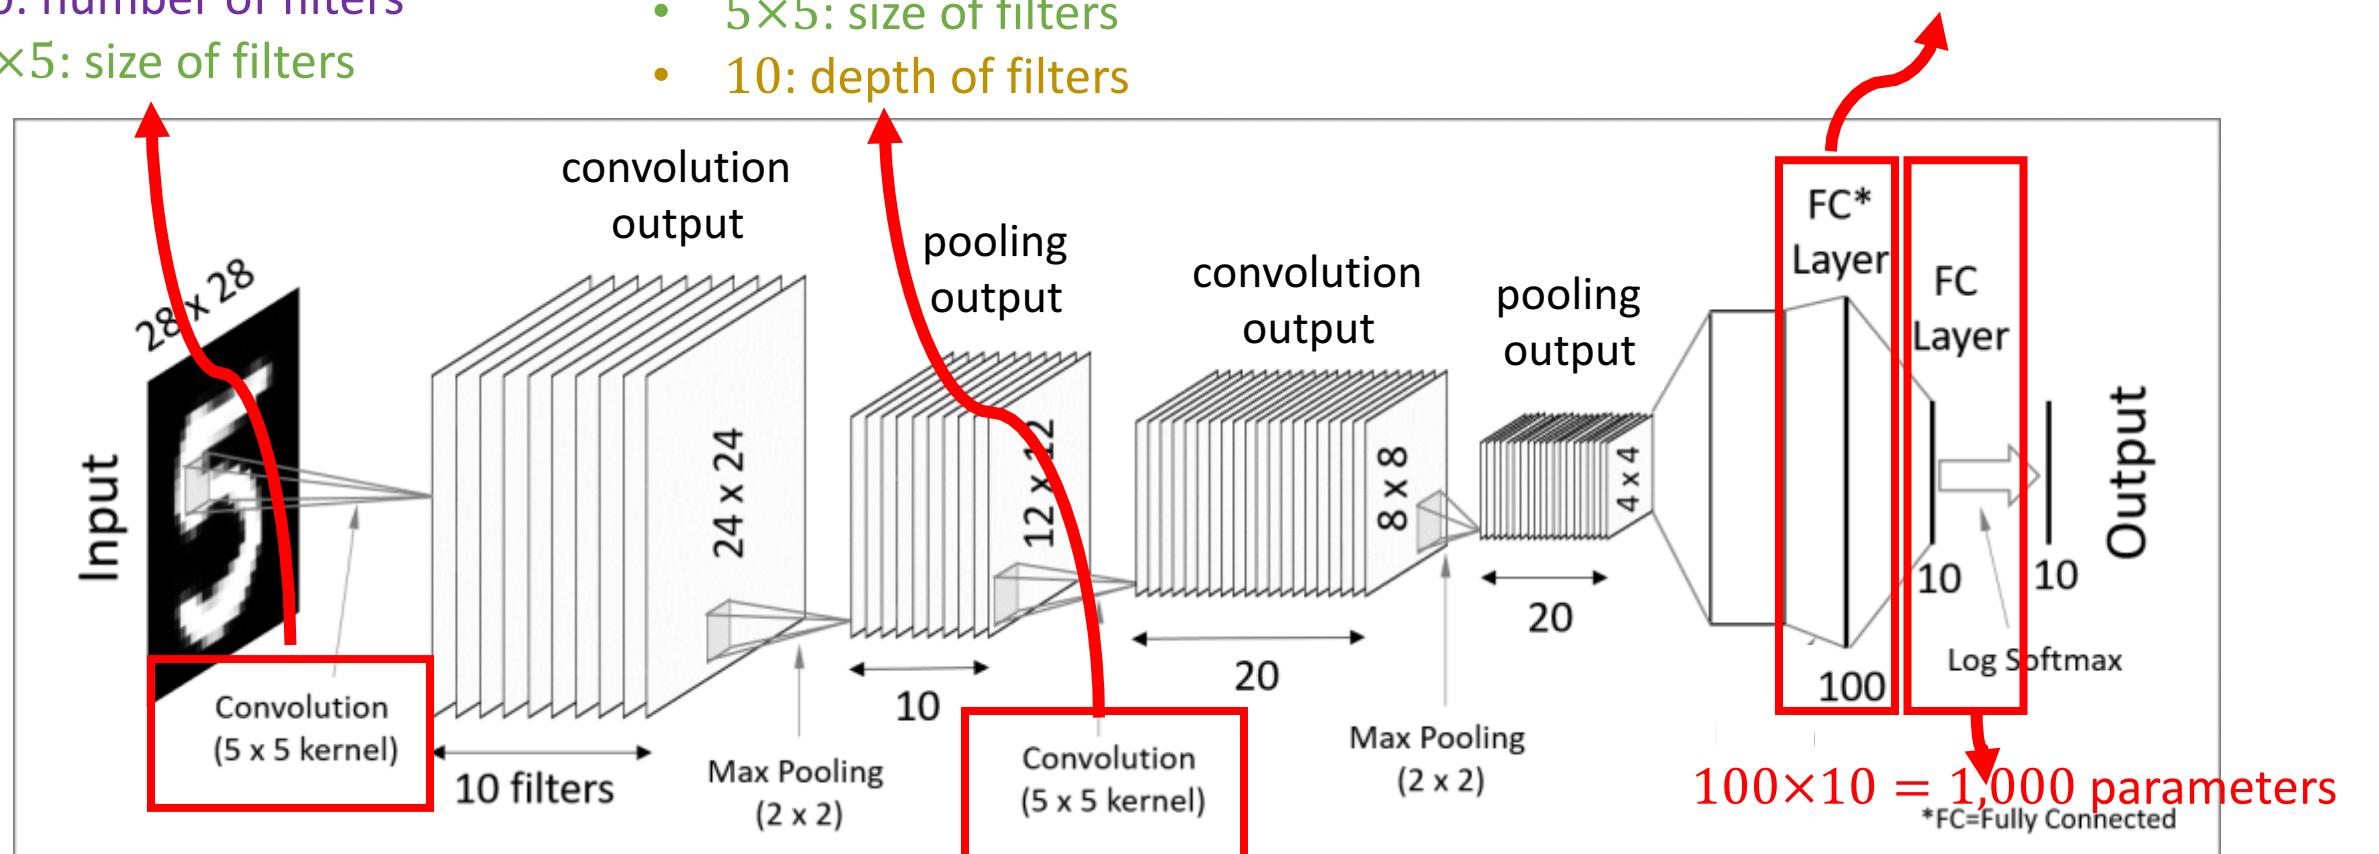
$$10 \times 5 \times 5 = 250 \text{ parameters}$$

- 10: number of filters
- 5×5: size of filters

$$20 \times 5 \times 5 \times 10 = 5000 \text{ parameters}$$

- 20: number of filters
- 5×5: size of filters
- 10: depth of filters

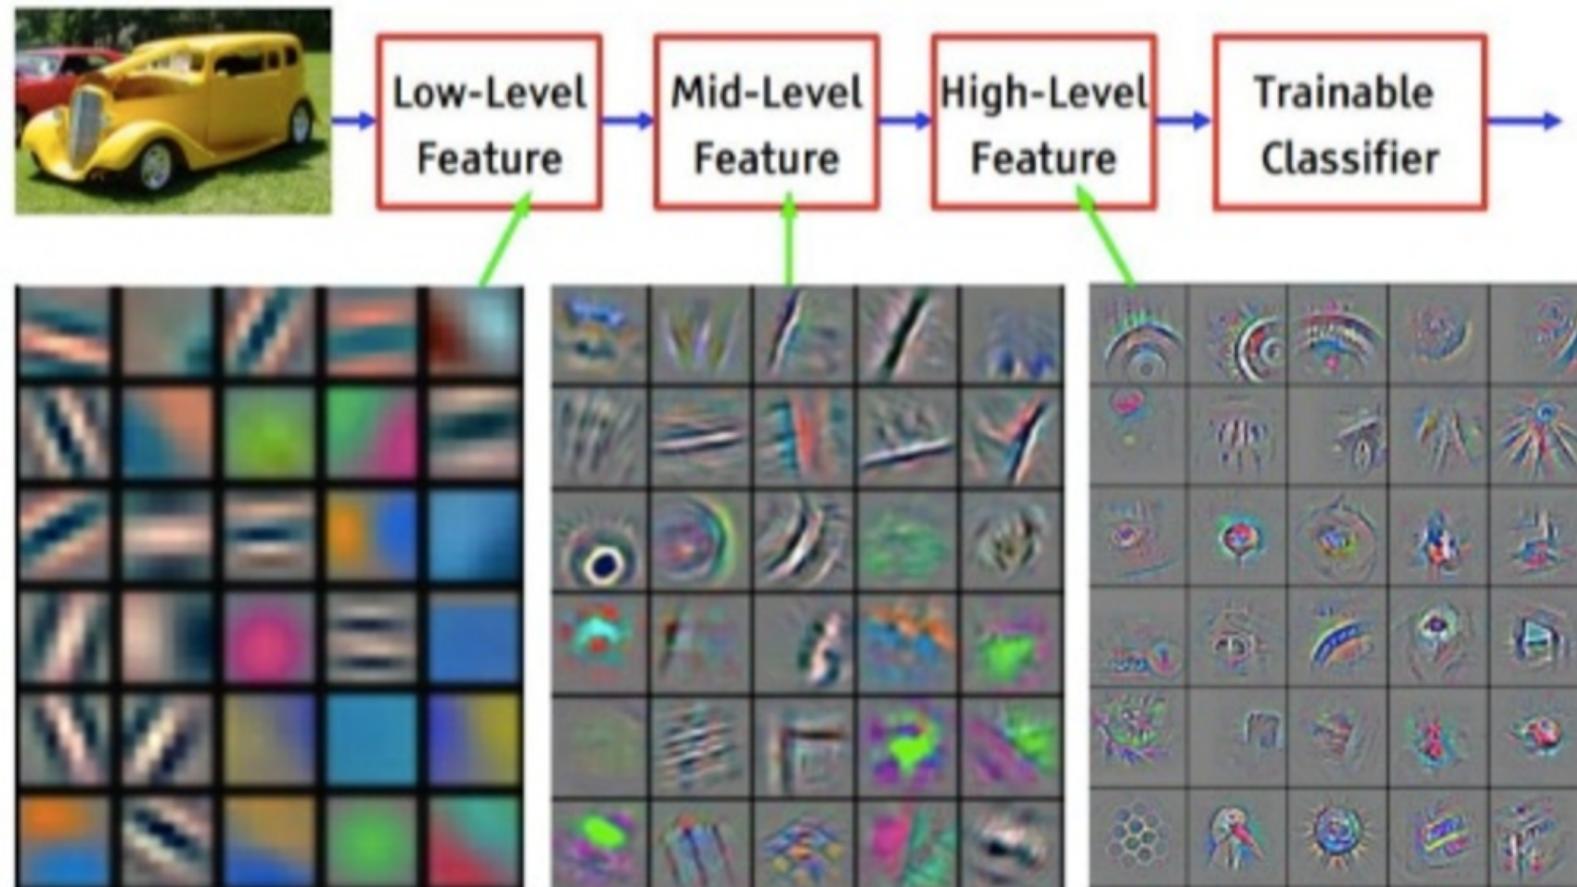
$$320 \times 100 = 32,000 \text{ parameters}$$



# CNN: Re-Cap

**Question:** Why convolutional layers?

**Answer:** detect features like edge, shape, pattern, etc.



# CNN: Re-Cap

**Question:** Why convolutional layers?

**Answer:** detect features like edge, shape, pattern, etc.

**Question:** Why activation functions like ReLU?

**Answer:** Add **nonlinearity** to increase the expressive power. (Note that convolution is a **linear** function.) 增加非线性，以增加表达能力。

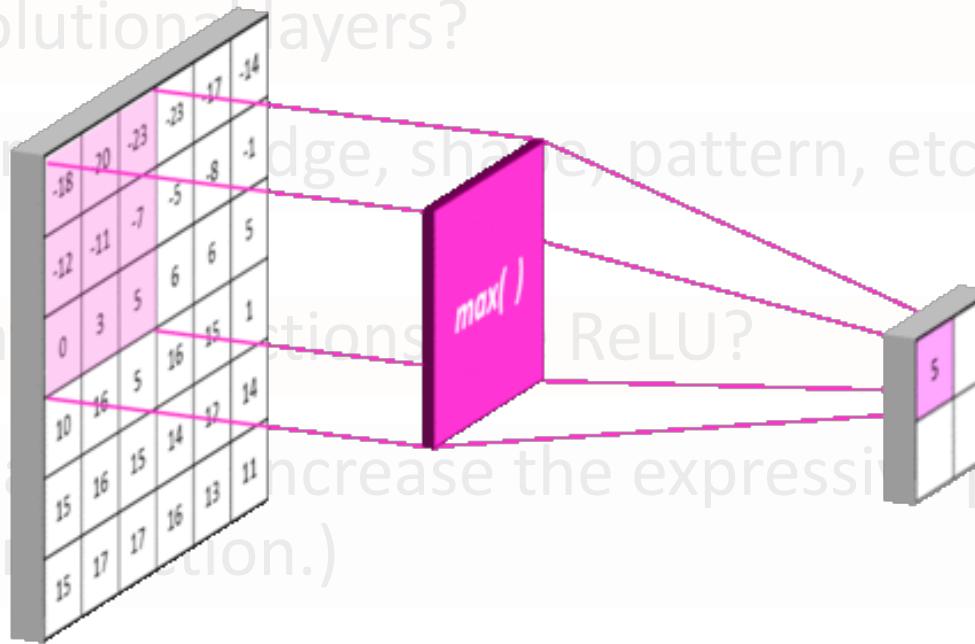
# CNN: Re-Cap

**Question:** Why convolutional layers?

**Answer:** detect features like edge, shape, pattern, etc.

**Question:** Why activation functions like ReLU?

**Answer:** Add nonlinearity to increase the expressiveness power. (Note that convolution is a linear operation.)



**Question:** Why pooling?

**Answer:** Drastically reduce the shape of features.

# Hyper-Parameters of CNN

# Hyper-Parameters for Network Structure

- Convolutional layers

# Filters

Filter shape

Stride

Zero-padding

- Pooling layers

MaxPool or AvgPool

Pool size

Pool stride

- Fully-connected layers

Width

- Activation functions

ReLU

SoftMax

Sigmoid

# Hyper-Parameters for Training

- Loss function.

Cross-entropy for classification

L1 or L2 for regression (the labels are continuous)

- Optimization algorithm (and its hyper-parameters, e.g., learning rate).

SGD

SGD with momentum

AdaGrad

RMSprop

- Random initialization.

Gaussian or uniform?

Scaling

# Implement a CNN Using Keras

# 1. Load and Process the MNIST Dataset

## Load data

```
from keras.datasets import mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()

print('Shape of x_train: ' + str(x_train.shape))
print('Shape of x_test: ' + str(x_test.shape))
print('Shape of y_train: ' + str(y_train.shape))
print('Shape of y_test: ' + str(y_test.shape))
```

Shape of x\_train: (60000, 28, 28)

Shape of x\_test: (10000, 28, 28)

Shape of y\_train: (60000,)

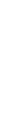
Shape of y\_test: (10000,)

# 1. Load and Process the MNIST Dataset

**Reshape:** convert the  $60000 \times 28 \times 28$  dataset to a  $60000 \times 28 \times 28 \times 1$  tensor.

```
x_train_vec = x_train.reshape((60000, 28, 28, 1)) / 255.0  
x_test_vec = x_test.reshape((10000, 28, 28, 1)) / 255.0  
  
print('Shape of x_train_vec is ' + str(x_train_vec.shape))
```

Shape of x\_train\_vec is (60000, 28, 28, 1)



单通道

# 1. Load and Process the MNIST Dataset

**One-hot encode:** convert the **labels** (scalars in  $\{0, 1, \dots, 9\}$ ) to 10-dim vectors.

```
def to_one_hot(labels, dimension=10):
    results = np.zeros((len(labels), dimension))
    for i, label in enumerate(labels):
        results[i, label] = 1.
    return results

y_train_vec = to_one_hot(y_train)
y_test_vec = to_one_hot(y_test)
print('Shape of y_train_vec is ' + str(y_train_vec.shape))
```

Shape of y\_train\_vec is (60000, 10)

# 1. Load and Process the MNIST Dataset

Partition the **training** set to **training** and **validation** sets.

```
# Partition to training and validation sets

rand_indices = np.random.permutation(60000)
train_indices = rand_indices[0:50000]
valid_indices = rand_indices[50000:60000]

x_valid_vec = x_train_vec[valid_indices, :, :, :]
y_valid_vec = y_train_vec[valid_indices, :]

x_train_vec = x_train_vec[train_indices, :, :, :]
y_train_vec = y_train_vec[train_indices, :]

print('Shape of x_valid_vec: ' + str(x_valid_vec.shape))
print('Shape of y_valid_vec: ' + str(y_valid_vec.shape))
print('Shape of x_train_vec: ' + str(x_train_vec.shape))
print('Shape of y_train_vec: ' + str(y_train_vec.shape))
```

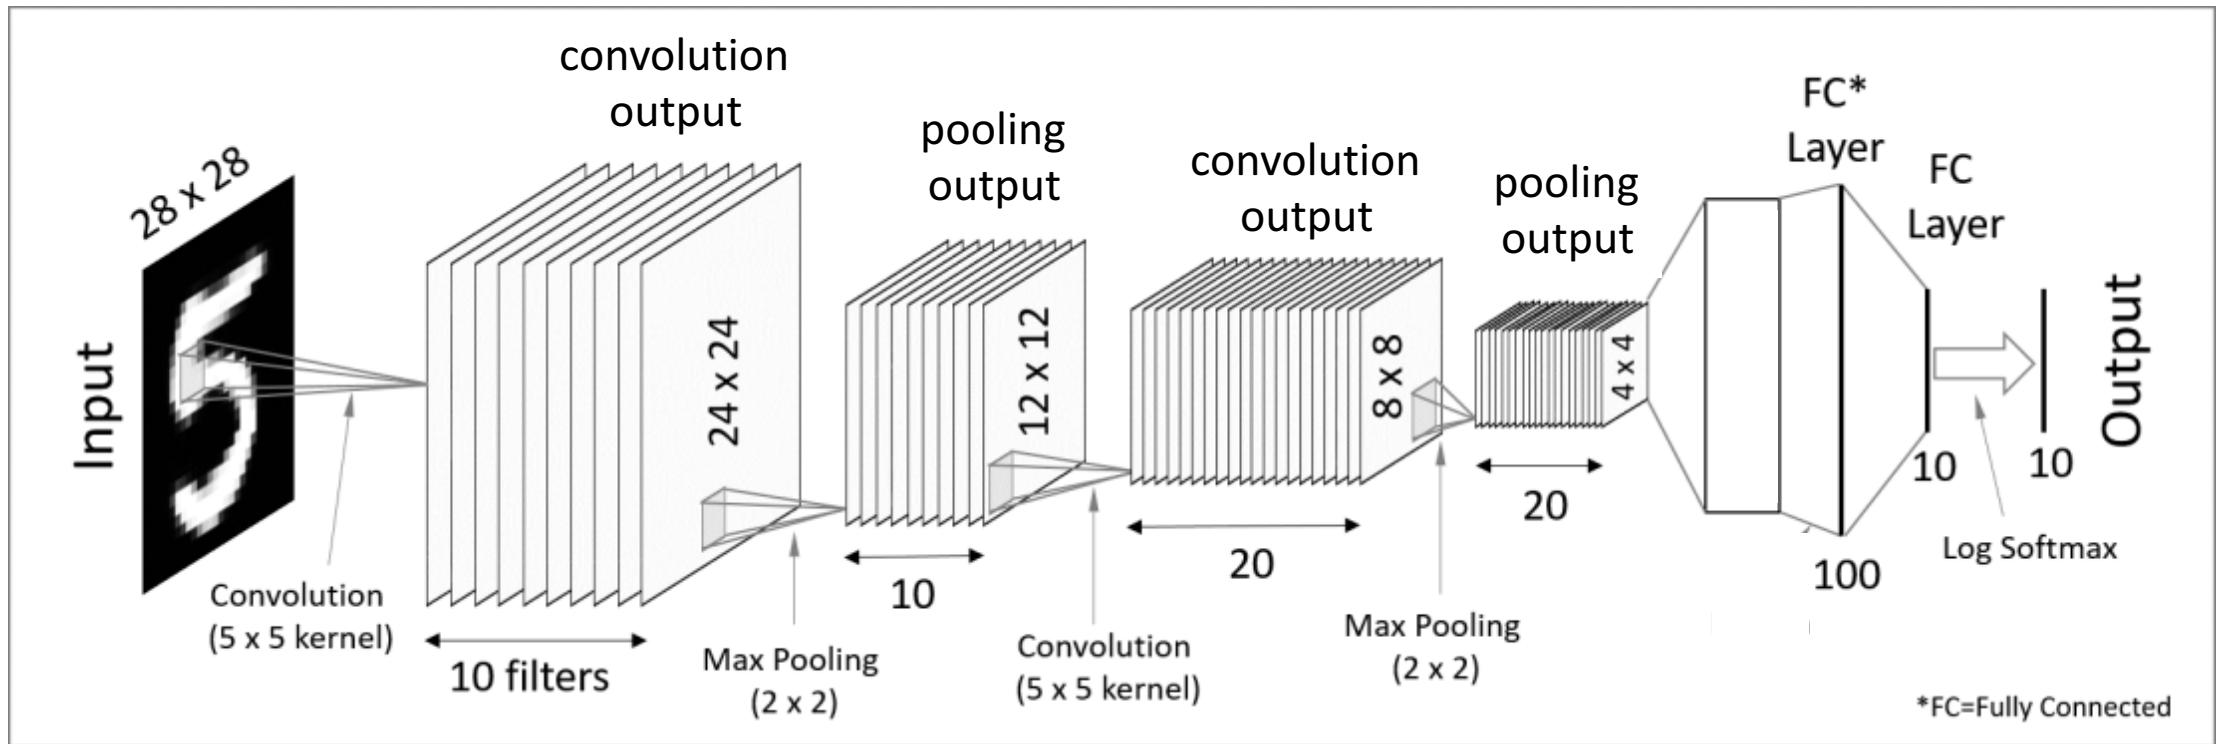
Shape of x\_valid\_vec: (10000, 28, 28, 1)

Shape of y\_valid\_vec: (10000, 10)

Shape of x\_train\_vec: (50000, 28, 28, 1)

Shape of y\_train\_vec: (50000, 10)

## 2. Build the CNN



## 2. Build the CNN

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(10, (5, 5), activation='relu', input_shape=(28, 28, 1)))
```

Filter number      Filter shape

- Default
- Stride: 1
  - Zero-padding: no padding

濾波器：(5, 5, 10)  
原图片：(28, 28, 1)

## 2. Build the CNN

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(10, (5, 5), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
```

pool\_size



Default

- Stride: pool\_size

## 2. Build the CNN

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(10, (5, 5), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2))) Output: 12×12×10
model.add(layers.Conv2D(20, (5, 5), activation='relu')) Output: 8×8×20
model.add(layers.MaxPooling2D((2, 2))) Output: 4×4×20
```

## 2. Build the CNN

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(10, (5, 5), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2))) Output: 12×12×10
model.add(layers.Conv2D(20, (5, 5), activation='relu')) Output: 8×8×20
model.add(layers.MaxPooling2D((2, 2))) Output: 4×4×20
model.add(layers.Flatten()) Output: 320
```

## 2. Build the CNN

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(10, (5, 5), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2))) Output: 12×12×10
model.add(layers.Conv2D(20, (5, 5), activation='relu')) Output: 8×8×20
model.add(layers.MaxPooling2D((2, 2))) Output: 4×4×20
model.add(layers.Flatten()) Output: 320
model.add(layers.Dense(100, activation='relu')) Output: 100
model.add(layers.Dense(10, activation='softmax')) Output: 10
```

## 2. Build the CNN

```
# print the summary of the model.  
model.summary()
```

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| =====                          |                    |         |
| conv2d_1 (Conv2D)              | (None, 24, 24, 10) | 260     |
| max_pooling2d_1 (MaxPooling2D) | (None, 12, 12, 10) | 0       |
| conv2d_2 (Conv2D)              | (None, 8, 8, 20)   | 5020    |
| max_pooling2d_2 (MaxPooling2D) | (None, 4, 4, 20)   | 0       |
| flatten_1 (Flatten)            | (None, 320)        | 0       |
| dense_1 (Dense)                | (None, 100)        | 32100   |
| dense_2 (Dense)                | (None, 10)         | 1010    |
| =====                          |                    |         |
| Total params: 38,390           |                    |         |
| Trainable params: 38,390       |                    |         |
| Non-trainable params: 0        |                    |         |

### 3. Train the CNN

Specify: optimization algorithm, learning rate (LR), loss function, and metric.

```
from keras import optimizers  
model.compile(optimizers.RMSprop(lr=0.0001),  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

### 3. Train the CNN

**Specify:** batch size and number of epochs.

```
history = model.fit(x_train_vec, y_train_vec,  
                     batch_size=128, epochs=50,  
                     validation_data=(x_valid_vec, y_valid_vec))
```

```
Train on 50000 samples, validate on 10000 samples  
Epoch 1/50  
50000/50000 [=====] - 12s 239us/step - loss: 1.2455 - acc: 0.7201 - val_loss: 0.4805 - val_acc: 0.8707  
Epoch 2/50  
50000/50000 [=====] - 12s 235us/step - loss: 0.3563 - acc: 0.8995 - val_loss: 0.2851 - val_acc: 0.9172  
Epoch 3/50  
50000/50000 [=====] - 12s 244us/step - loss: 0.2474 - acc: 0.9271 - val_loss: 0.2230 - val_acc: 0.9365
```

●  
●  
●

```
Epoch 49/50  
50000/50000 [=====] - 12s 240us/step - loss: 0.0244 - acc: 0.9924 - val_loss: 0.0439 - val_acc: 0.9875  
Epoch 50/50  
50000/50000 [=====] - 12s 239us/step - loss: 0.0238 - acc: 0.9927 - val_loss: 0.0446 - val_acc: 0.9881
```

# 4. Examine the Results

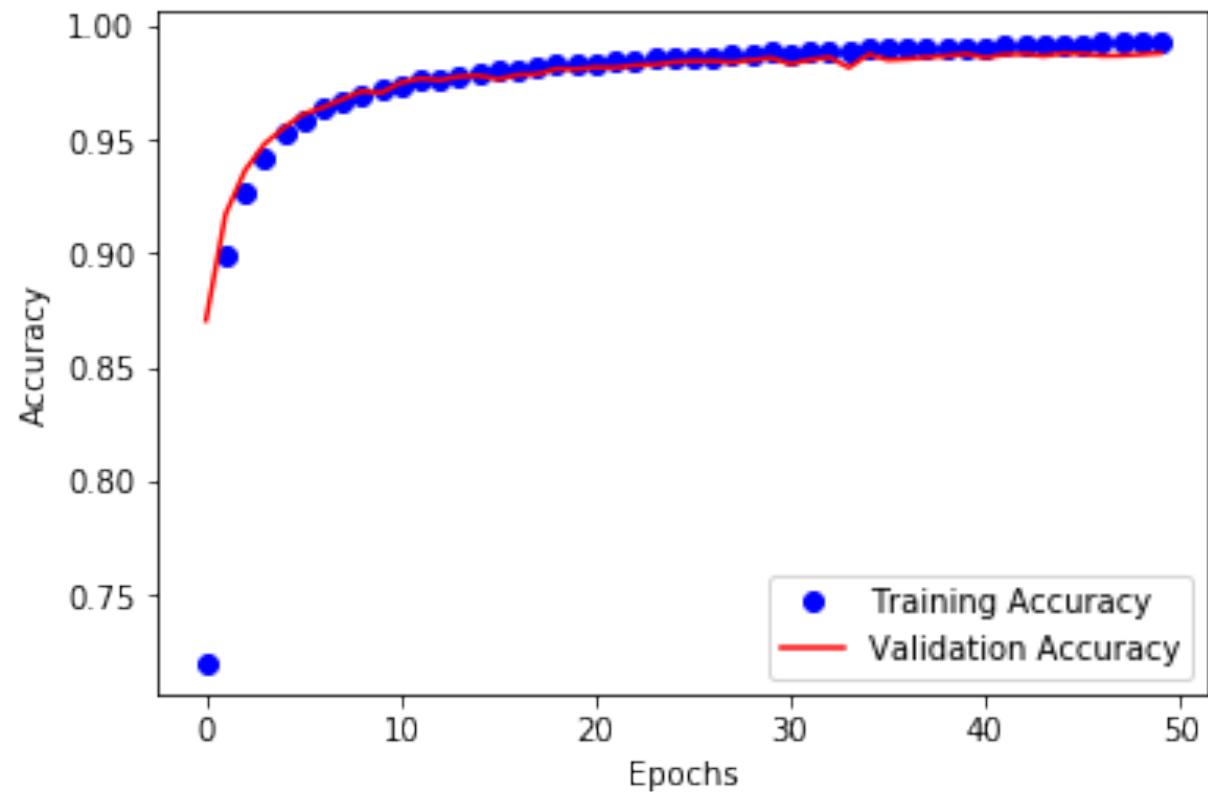
Plot the *accuracy* against *epochs* (1 epoch = 1 pass over the data).

```
import matplotlib.pyplot as plt
%matplotlib inline

epochs = range(50) # 50 is the number of epochs
train_acc = history.history['acc']
valid_acc = history.history['val_acc']
plt.plot(epochs, train_acc, 'bo', label='Training Accuracy')
plt.plot(epochs, valid_acc, 'r', label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

# 4. Examine the Results

Plot the *accuracy* against *epochs* (1 epoch = 1 pass over the data).



Why using the validation set?

- Tune the hyper-parameters, e.g., learning rate, filter shape, filter number, etc.
- Check overfitting or underfitting.

## 4. Examine the Results

Evaluate the model on the ***test*** set. (So far, the model has not seen the ***test*** set.)

```
loss_and_acc = model.evaluate(x_test_vec, y_test_vec)
print('loss = ' + str(loss_and_acc[0]))
print('accuracy = ' + str(loss_and_acc[1]))
```

```
10000/10000 [=====] - 1s 127us/step
loss = 0.03570800104729715
accuracy = 0.9883
```

## 4. Examine the Results

- Training accuracy: 99.3%
- Validation accuracy: 98.8%
- Test accuracy: 98.8%

# Summary

# 1. Convolution

- Matrix convolutional and tensor convolution.
- Concepts:
  - filter (kernel)
  - patch
  - output (feature map)
  - zero-padding
  - stride

## 2. Convolutional Neural Network

- Convolutional layers.
  - Filter number.
  - Filter shape.
  - Stride (default: one).
  - Zero-padding (default: no padding).

## 2. Convolutional Neural Network

- Convolutional layers.
  - Filter number.
  - Filter shape.
  - Stride (default: one).
  - Zero-padding (default: no padding).
- Activation functions.
- Pooling layers.
  - MaxPool, AveragePool, etc.
  - Pool size.
  - Pool stride (default: pool size).

### 3. Build and Train CNN using Keras

- Build a network.
  - Add Conv, Pool, FC layers.
- Compile the model.
  - Specify optimization algorithm, loss function, and metrics.
- Fit the model on training data.